
ML Unsupervised Techniques

Master Summary

Stefan Eder

Based on Hochreiter, *Machine Learning: Unsupervised Methods*, JKU Linz

This document is a compact summary of the course. It aims to cover all relevant concepts and results, but long derivations, detailed proofs, and worked examples are often condensed or omitted in favour of clarity and conciseness. Use it as a study aid alongside the lecture notes and slides, not as a replacement for them.

If you spot an error or a gap, please open an issue or pull request at github.com/stevetgit — corrections are always welcome.

Good study materials should be shared, not gatekept. If this helped you, take it as a reason to help the next colleague you see struggling with something.

Contents

1	Introduction	1
1.1	Machine Learning Introduction	1
1.2	Course Specific Introduction	1
1.3	Generative vs. Descriptive Models	3
2	Basic Terms and Concepts	4
2.1	Unsupervised Learning in Bioinformatics	4
2.2	Unsupervised Learning Categories	4
2.2.1	Generative Framework	4
2.2.2	Recoding Framework	5
2.2.3	Recoding and Generative Framework Unified	5
2.3	Quality of Parameter Estimation	5
2.4	Maximum Likelihood Estimator	6
2.5	Expectation Maximization	7
2.6	Maximum Entropy	8
3	Principal Component Analysis	10
3.1	The Method	10
3.2	Variance Maximization	10
3.3	Uniqueness	11
3.4	Properties of PCA	11
3.5	Examples	12
3.5.1	Iris Data Set	12
3.5.2	Multiple Tissue Data Set	12
3.6	Kernel Principal Component Analysis	13
3.6.1	The Method	13
4	Independent Component Analysis	14
4.1	Identifiability and Uniqueness	14
4.2	Measuring Independence	15
4.2.1	Mutual Information	15
4.2.2	Non-Gaussianity	15
4.3	Whitening and Rotation Algorithms	16
4.4	INFOMAX Algorithm	16
4.5	EASI Algorithm	17
4.6	FastICA Algorithm	17
4.7	ICA Extensions	18
4.8	ICA vs. PCA	18
4.9	Artificial ICA Examples	18
4.10	Real World ICA Examples	19
4.10.1	Iris Data Set	19
4.10.2	Multiple Tissue Data Set	19

4.11	Kurtosis Maximization Results in Independent Components	19
5	Factor Analysis	21
5.1	The Factor Analysis Model	21
5.2	Maximum Likelihood Factor Analysis	22
5.3	Factor Analysis vs. PCA and ICA	23
5.4	Artificial Factor Analysis Examples	23
5.5	Real World Factor Analysis Examples	23
5.5.1	Iris Data Set	23
5.5.2	Multiple Tissue Data Set	24
6	Scaling and Projection Methods	25
6.1	Projection Pursuit	25
6.2	Multidimensional Scaling	25
6.2.1	The Method	26
6.2.2	Examples	26
6.3	Non-negative Matrix Factorization	27
6.3.1	The Method	27
6.3.2	Examples	27
6.4	Locally Linear Embedding	28
6.4.1	The Method	28
6.4.2	Examples	29
6.5	Isomap	29
6.5.1	The Method	29
6.5.2	Examples	30
6.6	The Generative Topographic Mapping	30
6.6.1	The Method	30
6.6.2	Examples	31
6.7	t -Distributed Stochastic Neighbor Embedding	31
6.7.1	The Method	31
6.7.2	Examples	32
6.8	Self-Organizing Maps	32
6.8.1	The Method	32
6.8.2	Examples	33
7	Clustering	34
7.1	Mixture Models	34
7.1.1	The Method	34
7.1.2	Mixture of Gaussians	34
7.1.3	Mixture of Poissons	36
7.1.4	Examples	36
7.2	k -Means Clustering	37
7.2.1	The Method	37
7.2.2	Examples	38
7.3	Hierarchical Clustering	38
7.3.1	The Method	38
7.3.2	Examples	39
7.4	Similarity-Based Clustering	40
7.4.1	Aspect Model	40
7.4.2	Affinity Propagation	40
7.4.3	Similarity-based Mixture Models	42

8	Biclustering	45
8.1	Types of Biclusters	45
8.2	Overview of Biclustering Methods	45
8.3	FABIA Biclustering	46
8.4	Examples	47
9	Hidden Markov Models	49
9.1	Hidden Markov Models in Bioinformatics	49
9.2	Hidden Markov Model Basics	49
9.3	Expectation Maximization for HMM: Baum-Welch Algorithm	50
9.4	Viterbi Algorithm	50
9.5	Input Output Hidden Markov Models	50
9.6	Factorial Hidden Markov Models	50
9.7	Memory Input Output Factorial Hidden Markov Models	50
9.8	Tricks of the Trade	50
9.9	Profile Hidden Markov Models	50
10	Boltzmann Machines	51
10.1	The Boltzmann Machine	51
10.2	Learning in the Boltzmann Machine	51
10.3	The Restricted Boltzmann Machine	51

Introduction

1.1 Machine Learning Introduction

Machine learning is currently a major research topic in academia and industry alike. Applications span computer vision, speech recognition, recommender systems, and the analysis of large biological datasets. In the biomedical context, modern measurement technologies — mRNA concentrations measured via microarrays or next-generation sequencing, SNP arrays, copy number variation data — create a demand for methods that can extract structure from high-dimensional data without relying on labeled training examples.

The machine learning course series at JKU Linz comprises four complementary courses:

- **Basic Methods of Data Analysis:** summary statistics, PCA, ICA, multidimensional scaling, clustering, biclustering.
- **ML Supervised Methods:** classification and regression, SVMs, neural networks, deep learning, LSTM, ARMA/ARIMA.
- **ML Unsupervised Methods** (this course): MLE, MAP, maximum entropy, EM, PCA, ICA, factor analysis, mixture models, HMMs, Boltzmann machines.
- **Theoretical Concepts of Machine Learning:** estimation theory, VC-dimension, generalization bounds, convex optimization, Bayes theory.

The goal is not only to know individual methods, but to be able to *choose* the right approach for a given problem, understand the mathematical foundations well enough to adapt standard algorithms, and recognize when a model is likely to fail.

1.2 Course Specific Introduction

A learning problem is called *unsupervised* when the desired output is not given for the training data. In contrast to supervised learning, quality on unsupervised problems is measured as a *cumulative* quantity over all samples — there is no per-sample ground truth. Modern Generative AI (VAEs, GANs, diffusion models) is a prominent subtopic.

Objectives. Unsupervised methods are driven by different optimality criteria, depending on the problem:

- information content (maximizing information preserved after projection)
- orthogonality of extracted components
- statistical independence of components

- explained variance
- entropy of the representation
- likelihood: probability that the model produces the observed data
- distances within and between clusters

Use cases. Unsupervised methods are used to: explore data, find latent structure, visualize high-dimensional data in low dimensions, and compress representations. The unifying goal is to *understand and explore the data and generate new knowledge*.

Key concepts in this course. Maximum likelihood (ML), maximum a posteriori (MAP), maximum entropy, expectation maximization (EM), maximal variance, statistical independence, non-Gaussianity, sub- and super-Gaussian distributions, sparse codes and population codes.

Model selection and overfitting. The goal of learning is to select, from a model class, the model with the highest *generalization performance* — the model that best explains *future* data. Model selection, training, and learning are all the same thing. *Overfitting* occurs when the model is fitted to idiosyncratic training characteristics (noise, outliers, labeling errors) rather than to the true underlying structure; the model has high training performance but poor generalization.

Intuition & Mental Model

Underfitting (model too simple) and overfitting (model too complex) are the two failure modes. The goal is the best bias-variance trade-off: a model complex enough to capture the true structure but not so complex that it memorizes noise.

Parametric vs. nonparametric models. An important early decision is the *model class*:

- **Parametric models:** each parameter vector $\mathbf{w}$ represents one model. Examples: neural networks (synaptic weights), SVMs (weight vector $\mathbf{w}$), factor analysis. Learning corresponds to a path through parameter space. Disadvantages: different parameterizations can represent the same function; model complexity is implicitly controlled through the parameters.
- **Nonparametric models:** the model is locally constant or a superimposition of local pieces. Examples: k -nearest-neighbor (where k is a hyperparameter fixed before training), kernel density estimation, decision trees. The constant rules/kernels must be chosen a priori; hyperparameters (k , kernel width, splitting rules) are not learned from data.

Unsupervised methods covered in this course. Principal component analysis (PCA), independent component analysis (ICA), factor analysis, projection pursuit, k -means clustering, hierarchical clustering, mixture models (Gaussian, Poisson), self-organizing maps (SOMs), kernel density estimation, hidden Markov models (HMMs, factorial HMMs, IO-HMMs), Markov networks, restricted Boltzmann machines, auto-associators.

Methods fall into two broad families:

- **Projection methods:** produce a new representation of objects by down-projecting feature vectors into a lower-dimensional space while preserving neighborhoods or other desired properties. Examples: PCA (orthogonal components of maximal variance), ICA (statistically independent components), factor analysis (PCA with noise).
- **Generative models:** build an explicit probabilistic model of the observed data; the fitted model should match the observed data density.

1.3 Generative vs. Descriptive Models

Descriptive (Recoding) Model

A **descriptive model** provides an alternative representation or description of the data. It maps observations to a new space with desired properties (e.g. orthogonality, independence). Examples: PCA, ICA.

Generative Model

A **generative model** explicitly models the data generation process. It should produce the same distribution as the observed data: random source components $\mathbf{u}$ are drawn from a prior $p_{\mathbf{u}}$, passed through a model $g(\mathbf{u}; \mathbf{w})$, and the resulting distribution $p(\mathbf{x}; \mathbf{w})$ should match the empirical distribution $p(\mathbf{x})$. Examples: factor analysis, latent variable models, Boltzmann machines, HMMs.

Generative models allow the use of prior knowledge about the world and enable the prediction of new states of the data-generating process (e.g. predicting gene expression or brain activity). The natural training objective is maximum likelihood or MAP.

Descriptive models aim to find a richer or more interpretable representation rather than to model the generative process. Their objectives are properties of the representation: maximal variance (PCA), statistical independence (ICA), maximal non-Gaussianity (projection pursuit).

Exam trap

Parametric/nonparametric is *not* the same as generative/descriptive. Factor analysis is parametric *and* generative. k -NN is nonparametric *and* descriptive (no explicit density).

Basic Terms and Concepts

2.1 Unsupervised Learning in Bioinformatics

The primary driving application for many unsupervised methods is gene expression data. A microarray experiment produces a matrix of expression values: rows are genes (features), columns are samples (conditions, patients, time points). The data are high-dimensional (thousands of genes), noisy, and unlabeled.

Typical tasks include:

- finding groups of co-expressed genes (clustering, biclustering)
- identifying tissue types or disease subtypes (mixture models, ICA)
- visualizing the data in 2–3 dimensions (PCA, *t*-SNE, Isomap)
- detecting regulatory modules that activate in a subset of samples (FABIA)

A classical result is Spellman's cell-cycle data: PCA applied to yeast expression data over the cell cycle reveals the dominant temporal oscillation in the first two principal components, confirming the cyclic structure. Hierarchical clustering of microarray data yields dendrograms that group both genes and samples; the correlation coefficient bar on the dendrogram indicates relatedness of branches.

2.2 Unsupervised Learning Categories

All unsupervised methods fall into one of two frameworks:

Two Frameworks of Unsupervised Learning

- **Generative framework:** density estimation, hidden Markov models, mixture models. Objective: *maximum likelihood* or *maximum a posteriori*.
- **Recoding (descriptive) framework:** projection methods, PCA, ICA. Objective: *maximal variance*, orthogonality, independence, *maximum entropy*.

2.2.1 Generative Framework

In the generative framework, the world is modeled as a stochastic process: latent (hidden) sources $\mathbf{u}$ are drawn from a source distribution $p_{\mathbf{u}}$, passed through a world model $g(\mathbf{u}; \mathbf{w})$, and produce

observations $\mathbf{x}$. The learned model distribution $p(\mathbf{x}; \mathbf{w})$ must match the observed distribution $p(\mathbf{x})$.

Generative Model Objective

Find parameters $\mathbf{w}$ such that the model distribution

$$p(\mathbf{x}; \mathbf{w}) = \int_U p_u(\mathbf{u}) \delta(\mathbf{x} = g(\mathbf{u}; \mathbf{w})) d\mathbf{u}$$

is as close as possible to the empirical distribution $p(\mathbf{x})$. The loss function measures the distance between these two distributions; typically this is minimized via *maximum likelihood*.

Key examples of generative models: factor analysis, latent variable models, Boltzmann machines, hidden Markov models.

2.2.2 Recoding Framework

In the recoding (descriptive) framework, a model $g(\mathbf{x}; \mathbf{w})$ maps observations $\mathbf{x} \sim p(\mathbf{x})$ to projections $\mathbf{u}$. The resulting distribution $p(\mathbf{u}; \mathbf{w})$ should match a desired *target* distribution $p_t(\mathbf{u})$:

- **PCA:** projects to a low-dimensional space with maximal information conservation (maximal variance). Target: components are orthogonal and uncorrelated.
- **ICA:** projects into a space with statistically *independent* components (factorial code). Often characteristics of a factorial distribution are optimized: maximal entropy given fixed variance, or cumulants (kurtosis). Alternatively, the projection distribution is matched to a product of super-Gaussians.
- **Projection Pursuit:** projects such that components are maximally non-Gaussian (since Gaussian components arise from random projections of any distribution, by the CLT; non-Gaussian components are informative).

2.2.3 Recoding and Generative Framework Unified

Both frameworks can be seen as two sides of the same coin. In the generative direction: sources $\mathbf{u} \xrightarrow{g(\mathbf{u}; \mathbf{w})} \mathbf{x}$ (sources produce observations). In the recoding direction: observations $\mathbf{x} \xrightarrow{g^{-1}(\mathbf{x}; \mathbf{w})} \mathbf{u}$ (observations are encoded into sources). When the model g is invertible, these two views are mathematically equivalent and yield the same parameter estimates. This unification is exploited in ICA: the INFOMAX and likelihood approaches give identical learning rules for invertible linear models.

2.3 Quality of Parameter Estimation

Generative models must estimate true (but unknown) model parameters. Let the training data be $\{\mathbf{x}\} = \{\mathbf{x}^1, \dots, \mathbf{x}^l\}$, collected into the matrix $\mathbf{X} = (\mathbf{x}^1, \dots, \mathbf{x}^l)^\top \in \mathbb{R}^{l \times d}$. The true (unknown) parameter vector is $\mathbf{w}$; an estimator maps the data to $\hat{\mathbf{w}}$.

Estimator Quality Measures

$$\text{Unbiased: } \mathbb{E}_{\mathbf{X}}[\hat{\mathbf{w}}] = \mathbf{w}$$

$$\text{Bias: } b(\hat{\mathbf{w}}) = \mathbb{E}_{\mathbf{X}}[\hat{\mathbf{w}}] - \mathbf{w}$$

$$\text{Variance: } \text{Var}(\hat{\mathbf{w}}) = \mathbb{E}_{\mathbf{X}}[(\hat{\mathbf{w}} - \mathbb{E}_{\mathbf{X}}[\hat{\mathbf{w}}])^\top (\hat{\mathbf{w}} - \mathbb{E}_{\mathbf{X}}[\hat{\mathbf{w}}])]$$

$$\text{MSE: } \text{mse}(\hat{\mathbf{w}}) = \mathbb{E}_{\mathbf{X}}[(\hat{\mathbf{w}} - \mathbf{w})^\top (\hat{\mathbf{w}} - \mathbf{w})]$$

The fundamental bias-variance decomposition of the MSE is:

Bias-Variance Decomposition

$$\text{mse}(\hat{\mathbf{w}}) = \text{Var}(\hat{\mathbf{w}}) + b^2(\hat{\mathbf{w}})$$

Derivation sketch: Write $\hat{\mathbf{w}} - \mathbf{w} = (\hat{\mathbf{w}} - \mathbb{E}[\hat{\mathbf{w}}]) + (\mathbb{E}[\hat{\mathbf{w}}] - \mathbf{w})$, expand the squared norm, and note that the cross-term vanishes because $\mathbb{E}_{\mathbf{X}}[(\hat{\mathbf{w}} - \mathbb{E}[\hat{\mathbf{w}}])(\mathbb{E}[\hat{\mathbf{w}}] - \mathbf{w})^\top] = 0$ (since $\mathbb{E}[\hat{\mathbf{w}}] - \mathbf{w}$ is a constant w.r.t. the expectation).

The objective is to minimize the MSE. Note: the MSE in the parameter estimation sense (expected squared distance to true parameter) is different from the supervised regression loss (expected squared distance to true output).

Intuition & Mental Model

Unbiasedness and consistency are different properties. An *unbiased* estimator has $\mathbb{E}[\hat{\mathbf{w}}] = \mathbf{w}$ for *any* dataset size. A *consistent* estimator satisfies $\hat{\mathbf{w}} \rightarrow \mathbf{w}$ as $l \rightarrow \infty$ — its variance shrinks with more data, but it may be biased for finite l .

2.4 Maximum Likelihood Estimator

Likelihood

The **likelihood** of the dataset $\{\mathbf{x}\} = \{\mathbf{x}^1, \dots, \mathbf{x}^l\}$ under model parameters $\mathbf{w}$ is

$$\mathcal{L}(\{\mathbf{x}\}; \mathbf{w}) = p(\{\mathbf{x}\}; \mathbf{w}),$$

the probability that the model $p(\mathbf{x}; \mathbf{w})$ produces the observed data. For **iid** (independent, identically distributed) data this factorizes:

$$\mathcal{L}(\{\mathbf{x}\}; \mathbf{w}) = \prod_{i=1}^l p(\mathbf{x}^i; \mathbf{w}).$$

Working in log-space converts the product to a sum (negative log-likelihood, NLL):

$$-\ln \mathcal{L}(\{\mathbf{x}\}; \mathbf{w}) = -\sum_{i=1}^l \ln p(\mathbf{x}^i; \mathbf{w}).$$

Maximum likelihood estimation (MLE) is the dominant objective in unsupervised learning. It finds the parameters $\mathbf{w}$ that make the observed data most probable under the model:

$$\hat{\mathbf{w}}_{\text{ML}} = \arg \max_{\mathbf{w}} \mathcal{L}(\{\mathbf{x}\}; \mathbf{w}) = \arg \min_{\mathbf{w}} [-\ln \mathcal{L}(\{\mathbf{x}\}; \mathbf{w})].$$

Properties of the MLE

The maximum likelihood estimator is:

- **Invariant under reparameterization:** if $\hat{\mathbf{w}}$ is the MLE of $\mathbf{w}$, then $f(\hat{\mathbf{w}})$ is the MLE of $f(\mathbf{w})$ for any function f .
- **Asymptotically unbiased and efficient:** for large l , the MLE achieves the Cramér-Rao lower bound, i.e. it has the smallest possible variance among all unbiased estimators. It is asymptotically optimal.
- **Asymptotically consistent:** $\hat{\mathbf{w}} \rightarrow \mathbf{w}$ as $l \rightarrow \infty$.

MLE caveats

MLE can overfit for small datasets (it uses no prior information). For finite l , MLE may be biased: the classic example is the MLE of the variance of a Gaussian, $\hat{\sigma}^2 = \frac{1}{l} \sum (x_i - \bar{x})^2$, which underestimates the true variance (the unbiased estimator uses $l - 1$).

2.5 Expectation Maximization

Many generative models include *latent (hidden) variables* $\mathbf{u}$ that are not directly observed. The marginal likelihood then requires integrating over all latent configurations:

$$p(\mathbf{x}; \mathbf{w}) = \int_U p_{\mathbf{u}}(\mathbf{u}) \delta(\mathbf{x} = g(\mathbf{u}; \mathbf{w})) d\mathbf{u}.$$

This integral is in general intractable due to:

- hidden states (discrete or continuous latent variables),
- many-to-one output mappings (g is not injective),
- non-linearities in g .

The **Expectation Maximization (EM)** algorithm circumvents this by working with the *joint* distribution $p(\mathbf{x}, \mathbf{u}; \mathbf{w})$, which is easier to handle, and by constructing a tractable lower bound on the log-likelihood.

Derivation of the lower bound. Introduce an arbitrary distribution $Q(\mathbf{u} | \{\mathbf{x}\})$ over latent variables. Apply Jensen's inequality to the concave logarithm:

$$\begin{aligned} \ln \mathcal{L}(\{\mathbf{x}\}; \mathbf{w}) &= \ln \int_U p(\{\mathbf{x}\}, \mathbf{u}; \mathbf{w}) d\mathbf{u} \\ &= \ln \int_U Q(\mathbf{u} | \{\mathbf{x}\}) \frac{p(\{\mathbf{x}\}, \mathbf{u}; \mathbf{w})}{Q(\mathbf{u} | \{\mathbf{x}\})} d\mathbf{u} \\ &\geq \int_U Q(\mathbf{u} | \{\mathbf{x}\}) \ln \frac{p(\{\mathbf{x}\}, \mathbf{u}; \mathbf{w})}{Q(\mathbf{u} | \{\mathbf{x}\})} d\mathbf{u} := \mathcal{F}(Q, \mathbf{w}). \end{aligned}$$

$\mathcal{F}(Q, \mathbf{w})$ is the **free energy** (ELBO: evidence lower bound). Equivalently, the bound can be rewritten as

$$\mathcal{F}(Q, \mathbf{w}) = \ln \mathcal{L}(\{\mathbf{x}\}; \mathbf{w}) - D_{\text{KL}}(Q(\mathbf{u} | \{\mathbf{x}\}) \| p(\mathbf{u} | \{\mathbf{x}\}; \mathbf{w})),$$

which makes the E-step obvious: to maximize $\mathcal{F}$ over Q (with $\mathbf{w}$ fixed) one minimizes the KL divergence, which is zero iff $Q = p(\mathbf{u} | \{\mathbf{x}\}; \mathbf{w})$.

EM Algorithm

EM alternates between two steps until convergence:

$$\textbf{E-step: } Q_{k+1} = \arg \max_Q \mathcal{F}(Q, \mathbf{w}_k) \Rightarrow Q_{k+1} = p(\mathbf{u} \mid \{\mathbf{x}\}; \mathbf{w}_k)$$

$$\textbf{M-step: } \mathbf{w}_{k+1} = \arg \max_{\mathbf{w}} \mathcal{F}(Q_{k+1}, \mathbf{w})$$

Monotone convergence. Both steps can only increase (or leave unchanged) the lower bound $\mathcal{F}$. The E-step tightens the bound: at the beginning of the M-step, $\mathcal{F}(Q_{k+1}, \mathbf{w}_k) = \ln \mathcal{L}(\{\mathbf{x}\}; \mathbf{w}_k)$ (the bound is tight). The M-step then increases $\mathcal{F}$, which implies:

$$\ln \mathcal{L}(\{\mathbf{x}\}; \mathbf{w}_k) \leq \mathcal{F}(Q_{k+1}, \mathbf{w}_{k+1}) \leq \ln \mathcal{L}(\{\mathbf{x}\}; \mathbf{w}_{k+1}).$$

Therefore the log-likelihood is guaranteed to be non-decreasing at every EM iteration.

Intuition & Mental Model

The E-step does not change model parameters — it only updates the posterior distribution over latent variables. The M-step maximizes the expected complete-data log-likelihood, which for exponential family models has a closed-form solution. EM converges to a local maximum (not necessarily global).

EM is applied throughout this course: Baum-Welch algorithm for HMMs, fitting Gaussian mixtures, factor analysis, and (soft) ICA variants.

2.6 Maximum Entropy**Shannon Entropy (discrete case)**

For a discrete distribution $(p_1, p_2, \dots)$:

$$H = - \sum_{k \geq 1} p_k \log p_k, \quad \text{with the convention } 0 \cdot \log 0 = 0.$$

The **principle of maximum entropy** (Jaynes, 1957) states: among all distributions consistent with the available constraints (e.g. known moments), choose the one with *maximum entropy*. This distribution:

- makes the minimal assumptions beyond the stated constraints,
- corresponds to the most likely empirical distribution that could give rise to the observed features,
- connects statistical mechanics (thermodynamic equilibrium) and information theory.

Existence caveat

Not every class of distributions contains a maximum entropy member. For example, the class of all distributions with a given mean has arbitrarily large entropy (entropy can be pushed to ∞). The class of distributions with mean zero, second moment one, and third moment one has entropy bounded from above but the supremum is not attained.

Maximum entropy with moment constraints. Given moment constraints $\sum_{i=1}^n p(x_i) f_k(x_i) = F_k$ for $k = 1, \dots, m$ and normalization $\sum_{i=1}^n p(x_i) = 1$, the maximum entropy solution is the **Gibbs distribution**:

Gibbs Distribution (MaxEnt Solution)

$$p(x_i) = \frac{1}{Z(\lambda_1, \dots, \lambda_m)} \exp(\lambda_1 f_1(x_i) + \dots + \lambda_m f_m(x_i)),$$

with **partition function**

$$Z(\lambda_1, \dots, \lambda_m) = \sum_{i=1}^n \exp(\lambda_1 f_1(x_i) + \dots + \lambda_m f_m(x_i)).$$

The Lagrange multipliers λ_k are determined by the constraint equations:

$$F_k = \frac{\partial}{\partial \lambda_k} \log Z(\lambda_1, \dots, \lambda_m), \quad k = 1, \dots, m.$$

Standard MaxEnt examples

- **Normal distribution:** maximum entropy distribution given fixed mean and variance. Constraints: $\mathbb{E}[X] = \mu$ and $\mathbb{E}[X^2] = \mu^2 + \sigma^2$.
- **Uniform distribution:** maximum entropy on a bounded interval $[a, b]$. Only constraint: support.
- **Exponential distribution:** maximum entropy on $[0, \infty)$ given fixed mean. Constraint: $\mathbb{E}[X] = 1/\lambda$.

Intuition & Mental Model

The MaxEnt principle explains why Gaussian distributions appear so frequently in nature: they are the maximum entropy distributions for processes with finite variance. Any physical system evolving under energy constraints (finite second moment) will converge to a Gaussian-like distribution. In ICA, maximizing entropy of a nonlinear transform of the data is equivalent to maximizing likelihood of an independent-component model (INFOMAX).

Principal Component Analysis

3.1 The Method

Let the data matrix be $\mathbf{X} \in \mathbb{R}^{n \times m}$, where n is the number of samples and m is the number of features (the rows $\mathbf{x}_i \in \mathbb{R}^m$ are centered so that $\sum_i \mathbf{x}_i = \mathbf{0}$).

The **sample covariance matrix** is

$$\mathbf{C} = \frac{1}{n} \mathbf{X}^\top \mathbf{X} \in \mathbb{R}^{m \times m}.$$

Its eigendecomposition $\mathbf{C} = \frac{1}{n} \mathbf{U} \mathbf{D}_m \mathbf{U}^\top$ yields an orthogonal matrix $\mathbf{U}$ whose columns $\mathbf{u}_k$ are the principal directions (*loadings*), and a diagonal matrix $\mathbf{D}_m = \text{diag}(\lambda_1, \dots, \lambda_m)$ of eigenvalues sorted in decreasing order.

The **SVD** of $\mathbf{X}$ gives the same result: $\mathbf{X} = \mathbf{V} \mathbf{D} \mathbf{U}^\top$ with $\mathbf{V} \in \mathbb{R}^{n \times n}$, $\mathbf{D} \in \mathbb{R}^{n \times m}$ diagonal, and $\mathbf{U} \in \mathbb{R}^{m \times m}$ orthogonal. The eigenvalues of $\mathbf{C}$ are $\lambda_k = d_k^2/n$ where d_k are the singular values.

PCA Projection

The **principal component scores** (projections) are collected in

$$\mathbf{Y} = \mathbf{X} \mathbf{U} = \mathbf{V} \mathbf{D} \in \mathbb{R}^{n \times m}.$$

The k -th column $\mathbf{y}_k = \mathbf{X} \mathbf{u}_k$ is the projection of all data points onto the k -th principal direction. The original data can be reconstructed as the outer product sum

$$\mathbf{X} = \sum_{k=1}^m \mathbf{y}_k \mathbf{u}_k^\top.$$

Online learning: Oja's rule. A single-component PCA can be learned online (one sample at a time) using **Oja's rule**:

$$\mathbf{u}^{\text{new}} = \mathbf{u} + \eta(t \mathbf{x} - t^2 \mathbf{u}), \quad t = \mathbf{u}^\top \mathbf{x}.$$

This is a fixed-point iteration analogous to the power method.

3.2 Variance Maximization

PCA can be motivated as the projection that maximizes the variance of the projected data.

The first principal direction solves

$$\mathbf{u}_1 = \arg \max_{\|\mathbf{u}\|=1} \sum_{i=1}^n (\mathbf{u}^\top \mathbf{x}_i)^2 = \arg \max_{\|\mathbf{u}\|=1} \mathbf{u}^\top \mathbf{C} \mathbf{u}.$$

Variance Maximization Solution

Using the eigendecomposition $\mathbf{C} = \sum_k \lambda_k \mathbf{u}_k \mathbf{u}_k^\top$ and writing $\mathbf{u} = \sum_k a_k \mathbf{u}_k$, the objective becomes $\mathbf{u}^\top \mathbf{C} \mathbf{u} = \sum_k a_k^2 \lambda_k$ subject to $\sum_k a_k^2 = 1$. This is maximized by $a_1 = 1$, all others zero:

$$\mathbf{u}_1 = \text{eigenvector of } \mathbf{C} \text{ with largest eigenvalue } \lambda_1.$$

The k -th principal direction is the direction of maximal variance *orthogonal* to all previous $k - 1$ directions. The variance explained by the k -th component equals λ_k .

3.3 Uniqueness

PCA is **unique up to sign flips** of the principal directions, provided all eigenvalues are distinct. If two or more eigenvalues are equal (degenerate case), the corresponding eigenvectors span a subspace but their individual directions are arbitrary; any orthonormal basis of that subspace is equally valid.

Zero eigenvalues

Zero eigenvalues mean the data lies in a lower-dimensional affine subspace (rank-deficient covariance matrix). There can be *many* zero eigenvalues: if the data of dimension m spans only a k -dimensional subspace, then $m - k$ eigenvalues are zero. The corresponding principal directions can be discarded without any loss.

3.4 Properties of PCA

Let $\mathbf{y}^{(i)} = \mathbf{U}^\top \mathbf{x}_i$ be the projection of sample i .

Key Properties of PCA Projections

- **Zero mean:** $\sum_i y_k^{(i)} = 0$ for all k (because the data are centered).
- **Uncorrelated:** $\text{Cov}(y_k, y_l) = 0$ for $k \neq l$ (follows from orthogonality of $\mathbf{U}$).
- **Variance:** $\text{Var}(y_k) = \lambda_k$ (the k -th eigenvalue).
- **Optimal reconstruction:** truncating to the first l components gives the best rank- l approximation in terms of MSE:

$$\text{MSE}_l = \sum_{k=l+1}^m \lambda_k.$$

- **Fraction of variance** explained by the first l components: $\frac{\sum_{k=1}^l \lambda_k}{\sum_{k=1}^m \lambda_k}$.

Derivation of the MSE formula. Retain the first l components. The reconstruction of sample $\mathbf{x}_i$ is $\hat{\mathbf{x}}_i = \sum_{k=1}^l (\mathbf{u}_k^\top \mathbf{x}_i) \mathbf{u}_k$, so the residual is $\mathbf{x}_i - \hat{\mathbf{x}}_i = \sum_{k=l+1}^m (\mathbf{u}_k^\top \mathbf{x}_i) \mathbf{u}_k$. By orthonormality of the $\mathbf{u}_k$:

$$\|\mathbf{x}_i - \hat{\mathbf{x}}_i\|^2 = \sum_{k=l+1}^m (\mathbf{u}_k^\top \mathbf{x}_i)^2.$$

Averaging over all n samples and using $\frac{1}{n} \sum_{i=1}^n (\mathbf{u}_k^\top \mathbf{x}_i)^2 = \mathbf{u}_k^\top \mathbf{C} \mathbf{u}_k = \lambda_k$:

$$\text{MSE}_l = \frac{1}{n} \sum_{i=1}^n \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|^2 = \sum_{k=l+1}^m \mathbf{u}_k^\top \mathbf{C} \mathbf{u}_k = \sum_{k=l+1}^m \lambda_k.$$

Any other rank- l projection gives a larger MSE, so the top l eigenvectors are the optimal choice.

3.5 Examples

3.5.1 Iris Data Set

The Iris dataset has $n = 150$ samples (50 per species: *setosa*, *versicolor*, *virginica*) and $m = 4$ features (sepal length, sepal width, petal length, petal width).

PC1 accounts for approximately **92% of variance** and is the only component that separates the three species. Pairwise scatter plots of all four PCs show that only plots involving PC1 reveal a clear three-group structure; all other pairs (PC2 vs. PC3, etc.) show no meaningful separation. This tells us the discriminative signal is essentially one-dimensional.

3.5.2 Multiple Tissue Data Set

This microarray dataset contains expression values for 5,565 genes across 102 samples from four human tissue types: **breast** (Br), **prostate** (Pr), **lung** (Lu), and **colon** (Co).

Applied to all genes simultaneously:

- PC1 (5% variance) separates prostate from the rest.
- PC2 (3%) separates colon and also breast samples.
- PC3 (3%) partially separates lung.

Variance filtering. Selecting only the top-variance genes before PCA is justified for microarray data (most genes are essentially constant noise). Three filtered subsets were compared:

- **101 genes** (highest variance, XMultiF1): PC1 (36%) separates prostate; PC2 (14%) separates colon; PC3 (8%) separates breast and lung. Much cleaner separation than the full set.
- **13 genes** (XMultiF2, 57% PC1): similar tissue-level separation.
- **5 genes** (XMultiF3, 77% PC1): still separates prostate clearly but other tissues become hard to distinguish. Inspection reveals 4 of the 5 selected genes are highly correlated (ACPP, KLK2, MSMB, TRGC2 — all prostate-related; KRT5 is not).

Gene clustering. Applying hierarchical clustering with correlation as distance metric, then selecting the gene with maximal variance per cluster (**92 uncorrelated genes of maximal variance**) gives very similar results to variance-based filtering. The approach ensures genes are not redundant (correlated genes contribute the same information to PCA) while still selecting high-variance representatives.

3.6 Kernel Principal Component Analysis

3.6.1 The Method

Kernel PCA (KPCA) extends PCA to nonlinear projections by performing the linear operations of PCA inside a *reproducing kernel Hilbert space* (RKHS), to which the data are first mapped non-linearly:

$$\mathbf{x} \mapsto \Phi(\mathbf{x}).$$

Assuming the mapped data are centered ($\sum_i \Phi(\mathbf{x}_i) = \mathbf{0}$), the covariance in feature space is $\mathbf{C} = \frac{1}{n} \sum_i \Phi(\mathbf{x}_i) \Phi(\mathbf{x}_i)^\top$, and the Gram (kernel) matrix has entries $K_{ij} = \Phi(\mathbf{x}_j)^\top \Phi(\mathbf{x}_i) = k(\mathbf{x}_i, \mathbf{x}_j)$.

The eigenvalue problem $\mathbf{C}\mathbf{w} = \lambda\mathbf{w}$ cannot be solved in feature space directly (we only have K_{ij}), but the solution must lie in the span of the mapped data:

$$\mathbf{w} = \sum_{i=1}^n \alpha_i \Phi(\mathbf{x}_i).$$

Substituting and using K_{ij} yields the **kernel eigenvalue problem**:

$$n\lambda \mathbf{K}\boldsymbol{\alpha} = \mathbf{K}^2\boldsymbol{\alpha} \Rightarrow n\lambda \boldsymbol{\alpha} = \mathbf{K}\boldsymbol{\alpha}.$$

Normalization $\|\mathbf{w}\| = 1$ requires $\|\boldsymbol{\alpha}\| = 1/\sqrt{\tilde{\lambda}}$ where $\tilde{\lambda} = n\lambda$, so

$$\boldsymbol{\alpha} = \frac{\tilde{\boldsymbol{\alpha}}}{\sqrt{\tilde{\lambda}}}.$$

The projection of a new point $\mathbf{x}$ onto a kernel principal direction is then

$$\mathbf{w}^\top \Phi(\mathbf{x}) = \sum_{i=1}^n \alpha_i k(\mathbf{x}_i, \mathbf{x}).$$

Kernel PCA Algorithm

1. **Center** the Gram matrix: $\tilde{\mathbf{K}} = \mathbf{K} - \frac{1}{n}\mathbf{K}\mathbf{1}\mathbf{1}^\top - \frac{1}{n}\mathbf{1}\mathbf{1}^\top\mathbf{K} + \frac{1}{n^2}(\mathbf{1}^\top\mathbf{K}\mathbf{1})\mathbf{1}\mathbf{1}^\top$.
2. **Eigenvalues**: compute eigenvectors $\boldsymbol{\alpha}$ and eigenvalues $\tilde{\lambda}$ of $\tilde{\mathbf{K}}$.
3. **Normalize**: $\alpha_i^{\text{new}} = \alpha_i / (\|\boldsymbol{\alpha}\| \sqrt{n\lambda})$.
4. **Project** a new vector $\mathbf{x}$: center its kernel vector and compute $\mathbf{w}^\top \Phi(\mathbf{x}) = \sum_i \alpha_i k(\mathbf{x}_i, \mathbf{x})$.

Intuition & Mental Model

Kernel PCA can detect nonlinear structure that linear PCA misses. For example, concentric rings are linearly inseparable but perfectly separable in the RKHS of an RBF kernel. The price is that the representation depends on all training points (no compact weight vector $\mathbf{u}$), and choosing the right kernel function requires domain knowledge.

Independent Component Analysis

Independent component analysis (ICA) seeks a representation of the data as a linear mixture of **statistically independent** non-Gaussian sources.

ICA Model

Generative (mixing) model:

$$\mathbf{x} = \mathbf{U} \mathbf{y}, \quad p(\mathbf{y}) = \prod_{j=1}^l p(y_j).$$

$\mathbf{U} \in \mathbb{R}^{m \times l}$ is the mixing matrix; the sources y_j are statistically independent.

Descriptive (demixing) model:

$$\mathbf{y} = \mathbf{W} \mathbf{x}, \quad \mathbf{W} = \mathbf{U}^{-1} \quad (\text{for } l = m).$$

Matrix form: $\mathbf{X} = \mathbf{Y} \mathbf{U}^\top$, with $\mathbf{Y}^\top \mathbf{Y} = \mathbf{D}_m$ (decorrelated, and in fact statistically independent). Unlike PCA, $\mathbf{U}$ is **not required to be orthogonal**.

The canonical motivation is **blind source separation**: two speakers talk simultaneously; multiple microphones record linear mixtures of the two voices. ICA recovers the original signals from the mixtures without knowing the mixing matrix.

4.1 Identifiability and Uniqueness

The ICA solution is generally **not fully unique**: $\mathbf{x} = \mathbf{U} \mathbf{P}^{-1} \mathbf{P} \mathbf{y}$ gives another valid decomposition for any invertible $\mathbf{P}$. However, if $\mathbf{P}$ cannot mix the statistically independent components of $\mathbf{y}$, then $\mathbf{P}$ must be a product of a permutation matrix and a diagonal scaling matrix.

Darmois' Theorem (1953) — ICA Identifiability

Define $x_1 = \sum_j a_j y_j$ and $x_2 = \sum_j b_j y_j$ where y_j are independent random variables. If x_1 and x_2 are independent, then all y_j for which $a_j b_j \neq 0$ are *Gaussian*.

Consequence: for **non-Gaussian sources**, the ICA solution is unique up to permutation and scaling. The Gaussian case is the only exception.

ICA assumptions.

- Sources are **non-Gaussian** (at most one source can be Gaussian).
- $l \leq m$ (at least as many observations as sources).
- $\mathbf{U}$ has full rank l .

Source densities $p(y_i)$ are typically approximated by super-Gaussian or unimodal distributions (generative framework).

Two criteria for independence.

1. **Mutual information** between components.
2. **Non-Gaussianity** of components.

4.2 Measuring Independence

4.2.1 Mutual Information

The **mutual information** of a vector-valued random variable measures how far the joint distribution is from a factorial (product) distribution:

$$I(y_1, \dots, y_l) = \sum_{j=1}^l H(y_j) - H(\mathbf{y}),$$

where $H(\mathbf{a}) = -\int p(\mathbf{a}) \ln p(\mathbf{a}) d\mathbf{a}$ is the (differential) entropy. $I \geq 0$, with equality if and only if all components are statistically independent.

For the linear demixing $\mathbf{y} = \mathbf{W}\mathbf{x}$, using the change-of-variables formula $p(\mathbf{y}) = p(\mathbf{x})/|\mathbf{W}|$:

$$I(y_1, \dots, y_m) = \sum_{j=1}^m H(y_j) - H(\mathbf{x}) - \ln |\mathbf{W}|.$$

Since $H(\mathbf{x})$ is constant in $\mathbf{W}$, minimizing mutual information amounts to maximizing $\ln |\mathbf{W}| - \sum_j H(y_j)$, which is the ICA objective.

4.2.2 Non-Gaussianity

Negentropy. The **negentropy** measures how non-Gaussian a distribution is:

$$J(\mathbf{y}) = H(\mathbf{y}_{\text{gauss}}) - H(\mathbf{y}),$$

where $\mathbf{y}_{\text{gauss}}$ is a Gaussian with the same covariance as $\mathbf{y}$. $J \geq 0$ (Gaussian has maximum entropy for fixed variance); maximizing negentropy minimizes mutual information. Direct estimation of negentropy is difficult in practice.

Kurtosis. Non-Gaussianity can be quantified via cumulants, in particular the **kurtosis** (fourth cumulant):

$$\kappa_4 = \mathbb{E}(x^4) - 3(\mathbb{E}(x^2))^2.$$

For a Gaussian, $\kappa_3 = \kappa_4 = 0$.

- $\kappa_4 > 0$: **super-Gaussian** (heavier tails than Gaussian, e.g. Laplace distribution).
- $\kappa_4 < 0$: **sub-Gaussian** (lighter tails, e.g. uniform distribution).

Key properties:

- Additivity: $\kappa_4(x_1 + x_2) = \kappa_4(x_1) + \kappa_4(x_2)$ for independent x_1, x_2 .
- Scaling: $\kappa_4(\alpha x) = \alpha^4 \kappa_4(x)$.
- Mixing *reduces* kurtosis toward zero (towards Gaussian by CLT). Therefore **maximizing kurtosis** recovers the unmixed sources.

Kurtosis is not robust: as a fourth-order moment, it is heavily influenced by outliers.

Sparseness. Maximizing kurtosis is equivalent to maximizing **sparseness**: a sparse variable is almost always near zero but occasionally takes large values (relative to a Gaussian with the same variance). Sparseness does *not* imply small variance.

Contrast functions. Several contrast functions are used in practice to measure non-Gaussianity:

- Kurtosis: $\kappa_4(y)$.
- Mixed cumulants: $\frac{1}{12}\kappa_3^2(y) + \frac{1}{48}\kappa_4^2(y)$ (normalized to zero mean, unit variance).
- $|\mathbb{E}_y[G(y)] - \mathbb{E}_\nu[G(\nu)]|^p$ where ν is a standard Gaussian and G is a nonlinearity; robust choices include $G(x) = \log \cosh(ax)$ and $G(x) = \exp(-ax^2/2)$ with $a \geq 1$.

4.3 Whitening and Rotation Algorithms

ICA is not unique up to scaling, so the standard approach is to first remove the scaling ambiguity by **whitening** (also called sphering), then find an orthogonal **rotation** that maximizes non-Gaussianity.

Whitened data satisfies $\mathbf{Y}^\top \mathbf{Y} = \mathbf{I}$. Starting from $\mathbf{X} = \mathbf{Y}\mathbf{U}^\top$ with $\mathbf{C} = \frac{1}{n}\mathbf{X}^\top \mathbf{X}$:

$$\mathbf{I} = \mathbf{C}^{-1/2} \frac{1}{n} \mathbf{X}^\top \mathbf{X} \mathbf{C}^{-1/2} = \frac{1}{n} \mathbf{C}^{-1/2} \mathbf{U} \mathbf{Y}^\top \mathbf{Y} \mathbf{U}^\top \mathbf{C}^{-1/2}.$$

This means $\hat{\mathbf{U}} = \frac{1}{\sqrt{n}} \mathbf{C}^{-1/2} \mathbf{U}$ is orthogonal, and the ICA solution decomposes as:

$$\mathbf{Y} = \frac{1}{\sqrt{n}} \mathbf{X} \mathbf{C}^{-1/2} \hat{\mathbf{U}}.$$

The two-step procedure is:

1. **Whiten:** compute $\tilde{\mathbf{X}} = \mathbf{X} \mathbf{C}^{-1/2}$.
2. **Rotate:** find the orthogonal matrix $\hat{\mathbf{U}}$ that maximizes the contrast function (super-Gaussianity / sparseness).

4.4 INFOMAX Algorithm

INFOMAX (Bell & Sejnowski, 1995) minimizes mutual information by maximizing the entropy of a nonlinear transform of $\mathbf{y}$:

$$\mathbf{y} = \mathbf{W} \mathbf{x}, \quad H(\mathbf{g}(\mathbf{y})) \text{ is maximized,} \quad \mathbf{g}(\mathbf{y}) = (g(y_1), \dots, g(y_L))^\top.$$

At maximum entropy the mutual information $I(g(y_1), \dots, g(y_L)) = 0$, so the components are independent. The nonlinearity g is chosen to approximate the CDF of the source distribution; a common choice is $g(y_i) = \tanh(y_i)$ (super-Gaussian sources).

The entropy expands as

$$H(\mathbf{g}(\mathbf{y})) \approx H(\mathbf{x}) + \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^l |\ln g'(y_{ij})| + \ln |\mathbf{W}|.$$

INFOMAX Update Rules (Natural Gradient)

tanh ($g(y_j) = \tanh(y_j)$):

$$\Delta \mathbf{W} \propto (\mathbf{I} - 2\mathbf{g}(\mathbf{y})\mathbf{y}^\top)\mathbf{W}$$

sigmoid ($g(y_j) = 1/(1 + e^{-y_j})$):

$$\Delta \mathbf{W} \propto (\mathbf{I} + (1 - 2\mathbf{g}(\mathbf{y}))\mathbf{y}^\top)\mathbf{W}$$

INFOMAX is equivalent to a generative approach using **maximum likelihood**.

Intuition & Mental Model

The natural gradient (premultiplication by $\mathbf{W}^\top \mathbf{W}$) converts the standard gradient $(\mathbf{W}^\top)^{-1} \pm \dots \mathbf{x}^\top$ into the more elegant update involving $\mathbf{y} = \mathbf{W}\mathbf{x}$. It also makes the algorithm *equivariant*: performance does not depend on the scaling or conditioning of the input.

4.5 EASI Algorithm

EASI (Equivariant Adaptive Separation via Independence, Cardoso & Laheld, 1996) uses the update rule

$$\Delta \mathbf{W} \propto (\mathbf{I} - \mathbf{y}\mathbf{y}^\top - \mathbf{g}(\mathbf{y})\mathbf{y}^\top + \mathbf{y}\mathbf{g}^\top(\mathbf{y}))\mathbf{W},$$

where $\mathbf{g}$ is the same class of contrast functions as in INFOMAX (tanh or sigmoid). EASI is designed for online learning and is also equivariant by construction.

4.6 FastICA Algorithm

FastICA (Hyvärinen, 1999) is probably the most widely used ICA algorithm. It is a *fixed-point* iteration (analogous to Oja's rule for PCA) that maximizes non-Gaussianity as measured by a contrast function g with derivative g' :

FastICA Update

$$\mathbf{w}^{\text{new}} = \mathbb{E}[\mathbf{x} g(\mathbf{w}^\top \mathbf{x})] - \mathbb{E}[g'(\mathbf{w}^\top \mathbf{x})] \mathbf{w}.$$

The new $\mathbf{w}^{\text{new}}$ is then normalized to unit length.

Key properties:

- Whitening and rotation algorithm.
- Generalizes kurtosis maximization to arbitrary contrast functions g .
- Extended to extract multiple independent components simultaneously (deflation or symmetric extraction).

4.7 ICA Extensions

Beyond the basic square, invertible model several extensions exist:

- **Generative approach:** assumptions on the source densities $p(y_i)$, sub-Gaussian sources with specific priors.
- **Non-linear ICA:** mixing function $\mathbf{x} = g(\mathbf{y})$ is nonlinear; solutions are often not unique without additional constraints.
- **Overcomplete ICA:** more sources than observations ($l > m$); the system is underdetermined but useful for sparse representations of signals.
- **Undercomplete ICA:** fewer sources than observations ($l < m$); reduces to a lower-dimensional model.

4.8 ICA vs. PCA

ICA	PCA
Seeks <i>causes</i> of the data	Seeks <i>geometrical abstractions</i>
Statistically independent components	Decorrelated (orthogonal) components
Explains super-Gaussian structure	Explains all variance
Scale invariant	Not scale invariant
Unique up to scaling and permutation	Unique (up to sign)
Assumes super-Gaussian sources	No distributional assumptions
Components not ranked	Components ranked by eigenvalue

Intuition & Mental Model

PCA gives the best linear dimensionality reduction in terms of MSE; ICA finds the interpretation that is statistically most meaningful for independent source models. For Gaussian data both methods give the same result (any rotation of decorrelated Gaussians is equally independent), which is why ICA requires non-Gaussianity.

4.9 Artificial ICA Examples

Starting with 1,000 points from two independent **uniform distributions** (uniform = sub-Gaussian, $\kappa_4 < 0$), the three processing stages can be visualized:

1. **Original sources:** square-shaped scatter (axes-aligned).
2. **Mixing** ($\mathbf{X} = \mathbf{YU}^\top$): the square is rotated and sheared into a parallelogram/diamond shape — the components are now dependent.
3. **Whitening** ($\mathbf{XC}^{-1/2}$): the shape becomes circular (unit covariance); correlation is removed but statistical dependence remains.
4. **Rotation** (ICA): the circle is aligned back to a square, recovering the original independent sources.

Both INFOMAX and FastICA recover the original square (up to permutation and sign). For super-Gaussian sources (Laplace distribution), the shape is a four-pointed star, and the same pipeline recovers it correctly.

4.10 Real World ICA Examples

4.10.1 Iris Data Set

On the Iris data, ICs are **ordered by their impact on the observations** (column norms of the mixing matrix $\mathbf{U}$), *not* by variance. The first independent component explains approximately 90% of the variance in the data and likely encodes the *size of the blossom* (a single dominant biological factor). Unlike PCA, ICA does not guarantee components are ranked by variance.

4.10.2 Multiple Tissue Data Set

Applied to the multiple tissue microarray data (102 samples, 4 tissue types), the first three ICs provide biologically meaningful separations:

- **IC1** (29% variance): separates *secretory/reproductive* tissue (prostate and breast) from *internal organ* tissue (colon and lung) — a biologically interpretable axis.
- **IC2** (28%): separates prostate + lung from breast + colon.
- **IC3** (27%): separates prostate + colon from breast + lung.
- All combinations of the first three ICs yield clean separations except for partial overlap between breast and lung samples.

Genes most correlated to IC1 include SERPINA7, LAMB3, **AR**, CCNG2, KLF5, CCL20, SLC39A14, ATP1B1, GSTP1, LAD1. The **androgen receptor** (AR) is the 3rd most related gene to IC1 — consistent with IC1 separating reproductive/secretory tissues, since androgen signaling drives prostate gland differentiation and also plays a role in normal breast physiology.

Effect of increasing the number of components. With 8 or 20 ICA components, the overall tissue-level separation does *not* improve relative to 4 components. Instead, higher-order ICs focus on progressively finer-grained sub-populations (e.g. specific sub-groups within colon samples), and the main tissue axes (IC1: prostate vs. others; IC2: prostate+lung vs. breast+colon) shift in their index but remain. This illustrates a key difference from PCA: ICA components are not ordered by variance, and adding more components reveals biological sub-structure rather than explaining more total variance.

4.11 Kurtosis Maximization Results in Independent Components

The theoretical justification for kurtosis-based ICA rests on a key inequality: the kurtosis of a linear mixture of independent variables is always less than or equal to the *maximum* kurtosis of the individual sources. Formally, for $x = \mathbf{a}^\top \mathbf{y}$ with $\|\mathbf{a}\| = 1$ and y_j independent:

$$\kappa_4(x) = \sum_j a_j^4 \kappa_4(y_j) \leq \max_j |\kappa_4(y_j)|.$$

Equality holds when one $|a_j| = 1$ and all others are zero — i.e. when the projection aligns exactly with one source. Therefore **maximizing $|\kappa_4|$ over all unit-norm projections recovers a single source component**. Iterating (deflating each recovered component) separates all sources.

Limitations of kurtosis

Because kurtosis involves fourth moments, it is sensitive to outliers: a single extreme data point can dominate the estimate. More robust contrast functions such as $G(x) = \log \cosh(x)$ (used in FastICA) are preferred in practice when outliers are expected.

Factor Analysis

Factor analysis (FA) describes the **variability** of observations in terms of unobserved **latent variables** called *factors* and **noise**. The factors explain the correlations between observed variables; the remaining per-variable variance is explained by independent Gaussian noise. FA is a *generative* approach that models both the noise and the correlations explicitly.

5.1 The Factor Analysis Model

Factor Analysis Model

For centered data $\mathbf{x} \in \mathbb{R}^m$ with $l \leq m$ factors $\mathbf{z} \in \mathbb{R}^l$:

$$\mathbf{x} = \mathbf{\Lambda}\mathbf{z} + \boldsymbol{\varepsilon}, \quad \mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \quad \boldsymbol{\varepsilon} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Psi}),$$

where $\mathbf{\Lambda} \in \mathbb{R}^{m \times l}$ is the **factor loading matrix** and $\boldsymbol{\Psi} \in \mathbb{R}^{m \times m}$ is a **diagonal** noise covariance (so noise components are mutually independent). The model gives $\mathbf{x} \mid \mathbf{z} \sim \mathcal{N}(\mathbf{\Lambda}\mathbf{z}, \boldsymbol{\Psi})$.

Covariance decomposition. Marginalising over $\mathbf{z}$ gives $\mathbf{x} \sim \mathcal{N}(\mathbf{0}, \mathbf{\Lambda}\mathbf{\Lambda}^\top + \boldsymbol{\Psi})$. Empirically, with sample covariance $\mathbf{C} = \frac{1}{n}\mathbf{X}^\top\mathbf{X}$:

$$\mathbf{C} = \mathbf{\Lambda}\mathbf{\Lambda}^\top + \boldsymbol{\Psi}.$$

FA is therefore a **decomposition of the covariance matrix**: $\mathbf{\Lambda}$ captures the shared (dependent) variance; $\boldsymbol{\Psi}$ captures the unique (independent) variance per variable. Because $\boldsymbol{\Psi}$ is diagonal, all correlations between observations can *only* be explained by the factors.

In matrix form, $\mathbf{X} = \mathbf{Z}\mathbf{\Lambda}^\top + \mathbf{E}$ with model assumptions $\frac{1}{n}\mathbf{Z}^\top\mathbf{Z} = \mathbf{I}$, $\mathbf{Z}^\top\mathbf{E} = \mathbf{0}$, and $\frac{1}{n}\mathbf{E}^\top\mathbf{E} = \boldsymbol{\Psi}$.

Factor scores (estimation). Treating factor recovery as regression $\mathbf{Y} = \mathbf{X}\mathbf{A}$ and using the model approximations $\frac{1}{n}\mathbf{X}^\top\mathbf{X} \approx \mathbf{\Lambda}\mathbf{\Lambda}^\top + \boldsymbol{\Psi}$ and $\frac{1}{n}\mathbf{X}^\top\mathbf{Y} \approx \mathbf{\Lambda}$:

$$\hat{\mathbf{A}} = (\mathbf{\Lambda}\mathbf{\Lambda}^\top + \boldsymbol{\Psi})^{-1}\mathbf{\Lambda}, \quad \mathbf{Y} = \mathbf{X}(\mathbf{\Lambda}\mathbf{\Lambda}^\top + \boldsymbol{\Psi})^{-1}\mathbf{\Lambda}.$$

Applying the matrix inversion lemma to avoid inverting the $m \times m$ matrix:

$$\mathbf{Y} = \mathbf{X}\boldsymbol{\Psi}^{-1}\mathbf{\Lambda}(\mathbf{I} + \mathbf{\Lambda}^\top\boldsymbol{\Psi}^{-1}\mathbf{\Lambda})^{-1}.$$

Communality. The **communality** of observation variable x_j is the proportion of its variance explained by the factors:

$$c_j = \frac{\text{Var}(x_j) - \text{Var}(\varepsilon_j)}{\text{Var}(x_j)} = \frac{\sum_{k=1}^l \lambda_{jk}^2}{\Psi_{jj} + \sum_{k=1}^l \lambda_{jk}^2}.$$

Factor z_t individually contributes $\lambda_{jt}^2/(\Psi_{jj} + \sum_k \lambda_{jk}^2)$ to x_j . Like PCA, projecting onto l factors maximizes the variance that can be explained by l factors — but FA counts only signal variance, not noise.

Rotation non-uniqueness. The FA loading matrix is **not unique** up to orthogonal rotation: $\mathbf{\Lambda}\mathbf{z} = \mathbf{\Lambda}\mathbf{V}\mathbf{V}^\top\mathbf{z} = \mathbf{\Lambda}'\mathbf{z}'$ for any orthogonal $\mathbf{V}$. Rotations are therefore chosen to improve interpretability:

- **Varimax:** maximizes the squared loadings of each factor across all variables, so each factor has either large or small loadings — each variable is clearly associated with one factor.
- **Quartimax:** minimizes the number of factors needed per variable; each factor explains many variables but is usually hard to interpret.
- **Equimax:** a compromise between Varimax and Quartimax.

5.2 Maximum Likelihood Factor Analysis

Since $\mathbf{x} \sim \mathcal{N}(\mathbf{0}, \mathbf{\Lambda}\mathbf{\Lambda}^\top + \mathbf{\Psi})$, the log-likelihood over n samples is

$$\ln \mathcal{L}(\mathbf{\Lambda}, \mathbf{\Psi}) = \sum_{i=1}^n \left[-\frac{m}{2} \ln(2\pi) - \frac{1}{2} \ln |\mathbf{\Lambda}\mathbf{\Lambda}^\top + \mathbf{\Psi}| - \frac{1}{2} (\mathbf{x}^i)^\top (\mathbf{\Lambda}\mathbf{\Lambda}^\top + \mathbf{\Psi})^{-1} \mathbf{x}^i \right].$$

Maximising this has **no closed form**, so the **EM algorithm** is used with the latent factors $\mathbf{z}^i$ as the hidden variables.

FA EM Algorithm

E-step: compute the posterior of the factors for each sample (using the Gaussian conditional formula):

$$\mathbf{z}^i | \mathbf{x}^i \sim \mathcal{N}(\boldsymbol{\mu}_{\mathbf{z}^i | \mathbf{x}^i}, \boldsymbol{\Sigma}_{\mathbf{z}^i | \mathbf{x}^i}),$$

$$\boldsymbol{\mu}_{\mathbf{z}^i | \mathbf{x}^i} = (\mathbf{x}^i)^\top (\mathbf{\Lambda}\mathbf{\Lambda}^\top + \mathbf{\Psi})^{-1} \mathbf{\Lambda}, \quad \boldsymbol{\Sigma}_{\mathbf{z}^i | \mathbf{x}^i} = \mathbf{I} - \mathbf{\Lambda}^\top (\mathbf{\Lambda}\mathbf{\Lambda}^\top + \mathbf{\Psi})^{-1} \mathbf{\Lambda}.$$

Note: the posterior covariance $\boldsymbol{\Sigma}_{\mathbf{z}^i | \mathbf{x}^i}$ is the same for all data points (it does not depend on $\mathbf{x}^i$).

M-step: update the parameters by setting gradients to zero:

$$\mathbf{\Lambda}^{\text{new}} = \left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}^i \mathbb{E}_{\mathbf{z}^i | \mathbf{x}^i} [\mathbf{z}^i]^\top \right) \left(\frac{1}{n} \sum_{i=1}^n \mathbb{E}_{\mathbf{z}^i | \mathbf{x}^i} [\mathbf{z}^i (\mathbf{z}^i)^\top] \right)^{-1},$$

$$\mathbf{\Psi}^{\text{new}} = \frac{1}{n} \text{diag} \left(\sum_{i=1}^n \mathbf{x}^i (\mathbf{x}^i)^\top - \sum_{i=1}^n \mathbb{E}_{\mathbf{z}^i | \mathbf{x}^i} [\mathbf{z}^i] \mathbf{x}^i (\mathbf{\Lambda}^{\text{new}})^\top \right),$$

where $\mathbb{E}[\mathbf{z}^i (\mathbf{z}^i)^\top] = \boldsymbol{\mu}_{\mathbf{z}^i | \mathbf{x}^i} \boldsymbol{\mu}_{\mathbf{z}^i | \mathbf{x}^i}^\top + \boldsymbol{\Sigma}_{\mathbf{z}^i | \mathbf{x}^i}$.

Matrix inversion speedup. When $m \gg l$, directly inverting the $m \times m$ matrix $(\mathbf{\Lambda}\mathbf{\Lambda}^\top + \mathbf{\Psi})$ is expensive. The matrix inversion lemma gives an $l \times l$ inversion instead:

$$(\mathbf{\Lambda}\mathbf{\Lambda}^\top + \mathbf{\Psi})^{-1} = \mathbf{\Psi}^{-1} - \mathbf{\Psi}^{-1} \mathbf{\Lambda} (\mathbf{I} + \mathbf{\Lambda}^\top \mathbf{\Psi}^{-1} \mathbf{\Lambda})^{-1} \mathbf{\Lambda}^\top \mathbf{\Psi}^{-1}.$$

Since $\mathbf{\Psi}$ is diagonal, $\mathbf{\Psi}^{-1}$ is trivially cheap. Furthermore the sample covariance $\mathbf{C}$ only needs to be computed once per EM run.

5.3 Factor Analysis vs. PCA and ICA

	FA	PCA / ICA
Noise model	Explicit Gaussian noise Ψ	None (PCA: no noise; ICA: no noise)
Factors/sources	Gaussian priors $\mathcal{N}(\mathbf{0}, \mathbf{I})$	PCA: none; ICA: non-Gaussian
Objective	Max. signal variance (covariance fit)	PCA: all variance; ICA: independence
Uniqueness	Unique up to rotation	PCA: unique (up to sign); ICA: unique
Estimation	EM (no closed form)	PCA: eigen; ICA: gradient/fixed-point

Intuition & Mental Model

FA is the right choice when you believe observations contain genuine measurement noise that should be separated from the latent signal. PCA absorbs noise into the components (it explains *all* variance, noise included), while FA explicitly models it via Ψ . ICA additionally demands statistical independence of the sources and therefore requires non-Gaussian source distributions.

5.4 Artificial Factor Analysis Examples

To see when FA fails, consider 50-dimensional data generated from two independent **super-Gaussian** (Laplacian) latent sources. In 2-D projection the two sources form a characteristic “star” or diamond shape (uniform mass on four arms), which reflects their non-Gaussian character.

ICA correctly recovers this star shape: because the true sources are statistically independent and non-Gaussian, the ICA objective (contrast function or kurtosis) can identify them.

FA with two factors does *not* recover the star. FA’s generative model assumes the latent factors $\mathbf{z}$ are Gaussian; it can only model the global covariance structure $\mathbf{C} = \mathbf{\Lambda}\mathbf{\Lambda}^\top + \Psi$ and cannot distinguish the individual non-Gaussian sources. The FA scatter plots show a much worse demixing — the tissue subgroups blur together.

Why FA fails for non-Gaussian sources

FA is the maximum-likelihood model when factors are Gaussian. Forcing super-Gaussian data into a Gaussian-factor model distorts the latent structure: the EM algorithm converges, but to a solution that explains the observed covariance by smearing mass along incorrect axes. ICA is the right tool whenever the true sources are non-Gaussian.

5.5 Real World Factor Analysis Examples

5.5.1 Iris Data Set

Applying FA to the Iris data set with only **one factor** is sufficient to achieve a clean, linear separation of all three species (setosa, versicolor, virginica). The single maximum-likelihood factor captures the dominant axis of biological variation; a scatter plot of the factor scores colour-coded by species shows three nearly non-overlapping clusters arranged along a line.

5.5.2 Multiple Tissue Data Set

With **four factors** and no rotation the four components have distinct biological interpretations:

- **FA1** — cleanly separates prostate samples.
- **FA2** — partially separates breast from lung, but the overlap is substantial.
- **FA3** — separates colon samples.
- **FA4** — separates a subset of lung samples.

Applying **Varimax** rotation to the four factors changes the factor axes but does not change the overall fit. The pairwise scatter plots (all six pairs among FA1–FA4) show that *separation is slightly worse* than without rotation: the rotation mixes biological signal across components rather than concentrating it.

Quartimax rotation gives a result similar to Varimax — the separation quality is again slightly worse than the unrotated solution.

Rotation does not help here

Rotation methods such as Varimax and Quartimax search for a sparse, interpretable factor loading matrix. On this data the unrotated MLE already concentrates each factor on one tissue type, so rotation does not improve tissue separation and can even dilute it. Rotation is most valuable when the unrotated factors are hard to interpret biologically.

Scaling and Projection Methods

6.1 Projection Pursuit

All dimensionality-reduction methods in this chapter project data to a low-dimensional space, sometimes called *scaling* of the data. The goal is either to **visualise** the data or to produce a compact low-dimensional representation for downstream tasks, exploiting the fact that low-dimensional models are simpler and that noise and outliers are not represented in the projection.

Projection pursuit asks: among all linear projections $t = \mathbf{w}^\top \mathbf{x}$ (normalised to zero mean, unit variance), which direction $\mathbf{w}$ is most *interesting*? The classical answer is: the **least Gaussian** direction.

The rationale comes from information theory: among all distributions with a fixed mean and covariance, the Gaussian maximises entropy. Consequently, a projection that looks close to Gaussian is “boring” — it carries no structure beyond the second-order statistics already captured by PCA. A projection with low entropy (far from Gaussian) reveals genuine non-linear or non-Gaussian structure.

Projection Pursuit Objective

Find $\mathbf{w}$ that minimises the differential entropy $H(t)$ of the projected variable $t = \mathbf{w}^\top \mathbf{x}$, subject to $\text{Var}(t) = 1$.

This optimisation is difficult in practice because it requires estimating a density. Two qualitatively different types of structure can be found:

- **Unimodal super-Gaussians** — heavy-tailed, peaked distributions (single cluster with outliers).
- **Multimodal distributions** — several separated clusters.

In practice, projection pursuit uses the same contrast functions as ICA — kurtosis, negentropy approximations — as tractable surrogates for $H(t)$. Projection pursuit can therefore be seen as a precursor to ICA; ICA adds the constraint that multiple projected directions must be mutually independent.

6.2 Multidimensional Scaling

6.2.1 The Method

Multidimensional scaling (MDS) projects data to a low-dimensional target space $\mathbf{y}_i = f(\mathbf{x}_i; \mathbf{w})$ while keeping the *pairwise distances* between points as intact as possible.

Define:

$$\delta_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\| \quad (\text{original distances}), \quad d_{ij} = \|\mathbf{y}_i - \mathbf{y}_j\| \quad (\text{projected distances}).$$

The goal is $d \approx \delta$. Different **stress measures** formalise the mismatch:

$$R_1(d, \delta) = \frac{\sum_{i < j} (d_{ij} - \delta_{ij})^2}{\sum_{i < j} \delta_{ij}^2} \quad (\text{Kruskal's measure — penalises large absolute errors})$$

$$R_2(d, \delta) = \sum_{i < j} \left(\frac{d_{ij} - \delta_{ij}}{\delta_{ij}} \right)^2 \quad (\text{fractional/relative errors})$$

$$R_3(d, \delta) = \frac{1}{\sum_{i < j} \delta_{ij}} \sum_{i < j} \frac{(d_{ij} - \delta_{ij})^2}{\delta_{ij}} \quad (\text{Sammon mapping — compromise})$$

All three are minimised by gradient descent. The gradient of R_1 with respect to the projected position $\mathbf{y}_k$ is:

$$\frac{\partial}{\partial \mathbf{y}_k} R_1 = \frac{2}{\sum_{i < j} \delta_{ij}^2} \sum_{j \neq k} (d_{kj} - \delta_{kj}) \frac{\mathbf{y}_k - \mathbf{y}_j}{d_{kj}},$$

and analogous expressions hold for R_2 and R_3 (replacing the weight by $1/\delta_{kj}^2$ or $1/\delta_{kj}$, respectively). Viewing R as a potential energy, the gradients are **forces** that push or pull each point toward its correct distance from every other point.

A special case is **metric MDS** (also called *principal coordinates analysis*): when distances are Euclidean and one seeks a linear embedding, the optimal solution can be computed analytically via eigendecomposition of the centred distance matrix — no gradient descent is needed.

6.2.2 Examples

MDS is applied to the **Multiple Tissue** gene-expression data set (four tissue types: colon, lung, breast, prostate). The 101 or 13 genes with the largest variance across samples are selected as features.

- **Metric MDS** (101 features): colon samples (orange) cluster in the upper-left; prostate (green) separates to the right; lung (blue) and breast (red) form an overlapping cluster in the lower region. Using only 13 features gives a similar layout.
- **Kruskal's non-metric MDS** (101 and 13 features): the separation pattern is qualitatively the same as metric MDS.
- **Sammon's non-linear mapping** (101 features): the tissue groups are less well separated than with metric or Kruskal's measure. With only 13 features the clusters are *more spread out*, but the qualitative separation is still inferior to the other variants.

Overall, metric MDS and Kruskal's non-metric MDS perform similarly on this data set, and both outperform Sammon's mapping.

6.3 Non-negative Matrix Factorization

6.3.1 The Method

Non-negative matrix factorization (NMF) approximates a non-negative data matrix $\mathbf{X} \in \mathbb{R}^{n \times m}$ ($X_{ij} \geq 0$) as

$$\mathbf{X} \approx \mathbf{Y}\mathbf{U}^\top = \sum_{k=1}^l \mathbf{y}_k \mathbf{u}_k^\top,$$

where $\mathbf{Y} \in \mathbb{R}^{n \times l}$ and $\mathbf{U} \in \mathbb{R}^{m \times l}$ are also required to be non-negative ($Y_{ik} \geq 0$, $U_{jk} \geq 0$).

The non-negativity constraint forces a **parts-based representation**: every observation is reconstructed by *adding* non-negative parts, never by subtraction. This contrasts with PCA, whose eigenvectors can have both signs, producing holistic (whole-image) components.

Multiplicative update rules. Two objectives are in common use.

Objective 1 — KL divergence. For non-negative matrices normalised to sum to one:

$$D(\mathbf{A} \parallel \mathbf{B}) = \sum_{ij} \left(A_{ij} \log \frac{A_{ij}}{B_{ij}} + A_{ij} - B_{ij} \right).$$

Minimising $D(\mathbf{X} \parallel \mathbf{Y}\mathbf{U}^\top)$ leads to multiplicative updates:

$$Y_{ik} \leftarrow Y_{ik} \frac{\sum_{j=1}^m U_{jk} X_{ij} / (\mathbf{Y}\mathbf{U}^\top)_{ij}}{\sum_{j=1}^m U_{jk}}, \quad U_{jk} \leftarrow U_{jk} \frac{\sum_{i=1}^n Y_{ik} X_{ij} / (\mathbf{Y}\mathbf{U}^\top)_{ij}}{\sum_{i=1}^n Y_{ik}}.$$

Objective 2 — Frobenius norm. $\|\mathbf{X} - \mathbf{Y}\mathbf{U}^\top\|_F^2 = \sum_{ij} (X_{ij} - (\mathbf{Y}\mathbf{U}^\top)_{ij})^2$. Multiplicative updates are obtained by multiplying the fixed-point condition $\mathbf{X} = \mathbf{Y}\mathbf{U}^\top$ from the right by $\mathbf{U}$ (for $\mathbf{Y}$) or from the left by $\mathbf{Y}^\top$ (for $\mathbf{U}$):

$$Y_{ik} \leftarrow Y_{ik} \frac{(\mathbf{X}\mathbf{U})_{ik}}{(\mathbf{Y}\mathbf{U}^\top\mathbf{U})_{ik}}, \quad U_{jk} \leftarrow U_{jk} \frac{(\mathbf{Y}^\top\mathbf{X})_{kj}}{(\mathbf{Y}^\top\mathbf{Y}\mathbf{U}^\top)_{kj}}.$$

At a fixed point the left and right sides are equal.

Why multiplicative updates preserve non-negativity

Both update rules multiply each entry by a ratio of non-negative quantities. Starting from a non-negative initialisation, all entries remain non-negative throughout training.

NMF can be extended to **sparse NMF**, adding sparsity penalties on $\mathbf{Y}$ (few parts active per sample) and/or $\mathbf{U}$ (each part described by few features). In gene-expression analysis this corresponds to: part = pathway, few genes participate in each pathway, and only few pathways are active per sample.

6.3.2 Examples

Parts-based face representation.. Applied to a face database, NMF learns *local face parts* (eye regions, nose patches, mouth shapes). Vector quantisation (VQ) and PCA both learn holistic representations — entire faces or signed global components. The parts-based decomposition from NMF is more interpretable for many real-world objects.

Bicluster recovery (FABIA toy data).. A synthetic data matrix (1,000 genes, 100 samples, 13 biclusters) generated by the FABIA biclustering package provides a ground-truth test bed. The blocks of elevated expression are localised in specific gene/sample subsets; the background is near zero.

- **NMF (KL divergence, NMFDIV):** reconstructs the block structure well with small residual noise. The absolute factors and loadings matrices each show sparse, localised activity corresponding to individual biclusters.
- **NMF (Euclidean distance, NMFEU):** similarly good reconstruction; the residual error is small and structureless.
- **Sparse NMF (NMFSC):** enforcing sparseness misses some blocks entirely when the sparseness parameter is too large. The loadings matrix becomes nearly uniform across all genes (columns look identical) — the method cannot distinguish which genes belong to which bicluster. The takeaway is that *the sparseness parameter is difficult to tune*; too little gives the same result as plain NMF, too much destroys the bicluster structure.
- **FABIA biclustering:** the FABIA model (Chapter 8) recovers the bicluster structure more faithfully than any NMF variant, with sharper factor and loading matrices.

6.4 Locally Linear Embedding

6.4.1 The Method

Locally linear embedding (LLE) computes a low-dimensional, neighbourhood-preserving embedding via a nonlinear mapping. The key insight is that on a smooth manifold, each point can be approximated by a linear combination of its nearest neighbours; LLE finds the embedding that *preserves* those local linear relationships.

Step 1 — compute reconstruction weights.. For each point $\mathbf{x}_i$, find k nearest neighbours and solve for weights W_{ij} that best reconstruct $\mathbf{x}_i$ from its neighbours:

$$\varepsilon(\mathbf{W}) = \sum_i \left\| \mathbf{x}_i - \sum_{j=1}^k W_{ij} \mathbf{x}_j \right\|^2, \quad \text{subject to} \quad \sum_{j=1}^k W_{ij} = 1.$$

The constraint makes the weights invariant to rotations, rescalings, and translations of $\mathbf{x}_i$. Define the local covariance $C_{jk} = \mathbf{x}_j^\top \mathbf{x}_k$; the Lagrange multiplier enforcing the sum-to-one constraint is:

$$\lambda = \frac{1 - \sum_{jk} C_{jk}^{-1} \mathbf{x}_i^\top \mathbf{x}_k}{\sum_{jk} C_{jk}^{-1}},$$

and the optimal weights are:

$$W_{ij} = \sum_k C_{jk}^{-1} (\mathbf{x}_i^\top \mathbf{x}_k + \lambda).$$

Step 2 — compute the embedding.. Hold the weights W_{ij} fixed and find low-dimensional coordinates $\mathbf{y}_i$ that minimise the same reconstruction error:

$$\Phi(\mathbf{Y}) = \sum_{ij} M_{ij} \mathbf{y}_i^\top \mathbf{y}_j, \quad \text{where} \quad M_{ij} = \delta_{ij} - W_{ij} - W_{ji} + \sum_k W_{ki} W_{kj},$$

which can be written compactly as $\mathbf{M} = (\mathbf{I} - \mathbf{W})(\mathbf{I} - \mathbf{W})^\top$. Because the row sums of $\mathbf{W}$ are one, $(\mathbf{I} - \mathbf{W})^\top \mathbf{1} = \mathbf{0}$, so $\mathbf{M}$ always has a zero eigenvalue with eigenvector $\mathbf{1}$. The optimal d -dimensional embedding consists of the *bottom* d eigenvectors of $\mathbf{M}$ (those with the d smallest non-zero eigenvalues); the trivial zero eigenvector is discarded.

LLE implementation notes

- Neighbours can be found by k -NN (Euclidean), ε -ball, or KD trees for efficiency.
- If $k > m$ the local covariance $\mathbf{C}$ is rank-deficient; regularise with $\mathbf{C} \leftarrow \mathbf{C} + \varepsilon \mathbf{I}$ where $\varepsilon \approx 10^{-3} \text{tr}(\mathbf{C})$.
- The matrix $\mathbf{M}$ is sparse (only k off-diagonal entries per row), so its eigenvectors can be computed efficiently.

6.4.2 Examples

LLE succeeds on data that genuinely lives on a smooth low-dimensional manifold.

Swiss Roll. LLE “unrolls” the 3-D spiral manifold into a flat 2-D rectangle, recovering the original 2-D parameterisation of the roll.

Face images. Applied to a large collection of face photographs, LLE produces a 2-D embedding in which one axis encodes head pose and the other encodes expression or illumination — a smooth manifold structure in pixel space.

“S” curve. On a 3-D S-shaped surface LLE fails to produce a clean unrolling; the embedded scatter plot is disordered.

Multiple Tissue (gene expression). With $k \in \{5, 7, 9, 12\}$, LLE gives poor tissue separation in all cases. The four tissue types are not concentrated on a low-dimensional manifold in gene-expression space — they are spread throughout the high-dimensional space — so LLE’s assumption is violated.

6.5 Isomap**6.5.1 The Method**

Isomap computes a quasi-isometric low-dimensional embedding. Like LLE it is a non-linear projection, but instead of preserving local linear relationships it preserves **geodesic distances** — shortest paths along the data manifold.

Why geodesic rather than Euclidean?. For data lying on a curved manifold, the straight-line Euclidean distance between two far-apart points cuts through empty space and underestimates the true manifold distance. The geodesic distance follows the surface and is a better measure of “closeness” on the manifold.

Algorithm..

1. **Build the neighbourhood graph G .** Connect $\mathbf{x}_i$ and $\mathbf{x}_j$ if $d(\mathbf{x}_i, \mathbf{x}_j) < \varepsilon$ (ε -Isomap) or if i is among the k nearest neighbours of j (k -Isomap). Set each edge length to $d(\mathbf{x}_i, \mathbf{x}_j)$.
2. **Compute all-pairs shortest paths** using Floyd’s algorithm: initialise $d_G(\mathbf{x}_i, \mathbf{x}_j) = d(\mathbf{x}_i, \mathbf{x}_j)$ for linked pairs and ∞ otherwise, then iterate

$$d_G(\mathbf{x}_i, \mathbf{x}_j) \leftarrow \min(d_G(\mathbf{x}_i, \mathbf{x}_j), d_G(\mathbf{x}_i, \mathbf{x}_k) + d_G(\mathbf{x}_k, \mathbf{x}_j)).$$

Store the result as the geodesic distance matrix $\mathbf{D}_X$.

3. **Compute the embedding.** Convert distances to inner products via the doubly-centred transform

$$\tau(\mathbf{D}) = -\frac{1}{2} \mathbf{H} \mathbf{D}^2 \mathbf{H}, \quad \mathbf{H} = \mathbf{I}_n - \frac{1}{n} \mathbf{1}\mathbf{1}^\top,$$

where $\mathbf{D}^2$ is the element-wise square of $\mathbf{D}$. The embedding coordinates are given by $y_{ij} = \sqrt{\lambda_p} v_{ji}$, where λ_p and v_p are the p -th largest eigenvalue and eigenvector of $\tau(\mathbf{D}_X)$.

The objective that is minimised is $E = \|\tau(\mathbf{D}_X) - \tau(\mathbf{D}_Y)\|_{L^2}$, where $\mathbf{D}_Y$ collects Euclidean distances in the projected space. Setting the embedding to the top- l eigenvectors of $\tau(\mathbf{D}_X)$ is the exact analytic solution.

6.5.2 Examples

Hand images. $n = 2,000$ images of a hand opening and closing at varying wrist orientations ($k = 6$). Isomap recovers the underlying two-dimensional manifold: one axis encodes *wrist rotation*, the other *finger extension* — a semantically meaningful 2-D embedding of a high-dimensional image space.

Tree-count data (Barro Colorado Island). 50 one-hectare plots, 225 tree species. Classical metric MDS (cmdscale), k -Isomap ($k = 3, k = 5$), and ε -Isomap ($\varepsilon = 0.45$) give qualitatively different layouts of the 50 plots in 2-D. The neighbourhood parameter strongly influences the result.

Multiple Tissue. Results are not as good as with linear methods (metric MDS, FA) for any setting of ε . As with LLE, the gene-expression data do not lie on a smooth manifold, so geodesic-distance-based methods gain nothing over Euclidean methods and may even perform worse.

6.6 The Generative Topographic Mapping

6.6.1 The Method

The **Generative Topographic Mapping** (GTM) is a probabilistic, non-linear latent variable model designed as a principled alternative to Self-Organising Maps (SOMs). Like FA, GTM maps from a low-dimensional latent space to the observation space; unlike FA the mapping $\mathbf{x}(\mathbf{y}; \mathbf{w})$ is nonlinear.

Latent variables $\mathbf{y} \in \mathbb{R}^l$ are mapped to a manifold $\mathcal{S}$ embedded in $\mathbb{R}^m$ ($m > l$). Because a deterministic map from an l -dimensional space to an m -dimensional space concentrates all probability mass on a zero-volume surface (probability densities vanish), GTM places a **Gaussian ball** around each mapped point:

$$p(\mathbf{t} \mid \mathbf{y}, \mathbf{w}, \beta) = \left(\frac{\beta}{2\pi} \right)^{m/2} \exp\left(-\frac{\beta}{2} \|\mathbf{x}(\mathbf{y}, \mathbf{w}) - \mathbf{t}\|^2 \right).$$

The marginal density over observations is obtained by integrating over the latent space:

$$p(\mathbf{t} \mid \mathbf{w}, \beta) = \int p(\mathbf{t} \mid \mathbf{x}, \mathbf{w}, \beta) p(\mathbf{x}) d\mathbf{x}.$$

In practice the latent prior $p(\mathbf{y})$ is approximated by a uniform discrete distribution over L grid points $\{\mathbf{y}_j\}$:

$$p(\mathbf{y}) = \frac{1}{L} \sum_{j=1}^L \delta(\mathbf{y} - \mathbf{y}_j), \quad p(\mathbf{t} \mid \mathbf{w}, \beta) = \frac{1}{L} \sum_{j=1}^L p(\mathbf{t} \mid \mathbf{y}_j, \mathbf{w}, \beta).$$

This is a **constrained Gaussian mixture model** in $\mathbb{R}^m$: the L mixture centres are not free but must all lie on the manifold $\mathcal{S} = \{\mathbf{x}(\mathbf{y}_j; \mathbf{w})\}$.

The model is fitted by maximising the log-likelihood $\log \mathcal{L} = \sum_{i=1}^n \ln p(\mathbf{t}_i | \mathbf{w}, \beta)$ using EM.

6.6.2 Examples

A toy problem illustrates GTM: data are generated from a 1-D curve embedded in 2-D. The L latent grid points (shown as + in the initial configuration) are initialised along a straight line and each surrounded by its Gaussian noise distribution. After fitting, the latent points align with the data curve (right panel): the model has learned the non-linear manifold, and each observation is explained by the nearest Gaussian centre on the fitted curve.

6.7 t -Distributed Stochastic Neighbor Embedding

6.7.1 The Method

t -SNE maps each high-dimensional point to a 2- or 3-dimensional position such that similar points land nearby and dissimilar points land far apart.

Stochastic Neighbor Embedding (SNE).. For each pair in the original space, define the *conditional probability* that $\mathbf{x}_i$ would pick $\mathbf{x}_j$ as its neighbour:

$$p_{j|i} = \frac{\exp(-\|\mathbf{x}_i - \mathbf{x}_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|\mathbf{x}_i - \mathbf{x}_k\|^2 / 2\sigma_i^2)}.$$

The bandwidths σ_i are set so that the *perplexity*

$$\text{Perp}(P_i) = 2^{H(P_i)}, \quad H(P_i) = - \sum_{j \neq i} p_{j|i} \log_2 p_{j|i},$$

equals a user-chosen constant (the perplexity hyperparameter, typically 5–50). Here $H(P_i)$ is the Shannon entropy of the conditional distribution $P_i = \{p_{j|i}\}$ over the observation-space neighbours of i . A larger perplexity considers more neighbours as relevant. In the low-dimensional space define analogously:

$$q_{j|i} = \frac{\exp(-\|\mathbf{y}_i - \mathbf{y}_j\|^2)}{\sum_{k \neq i} \exp(-\|\mathbf{y}_i - \mathbf{y}_k\|^2)}.$$

The embedding is found by minimising the KL divergence $KL(P||Q) = \sum_{i \neq j} p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$ via gradient descent.

Problems with SNE — and how t -SNE fixes them..

1. **Difficult optimisation.** t -SNE *symmetrises* the joint distribution: $p_{ij} = (p_{j|i} + p_{i|j}) / (2n)$, giving simpler gradients.
2. **Crowding problem.** In 10 dimensions, 11 points can be mutually equidistant; there is no faithful 2-D map for them. t -SNE uses a **heavy-tailed Student's t -distribution** for the low-dimensional similarities:

$$q_{ij} = \frac{(1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|\mathbf{y}_k - \mathbf{y}_l\|^2)^{-1}}.$$

The heavier tail allows moderately distant points in the original space to be pushed further apart in 2-D without penalty, alleviating crowding.

Why the t -distribution helps

A Gaussian distribution has thin tails: forcing points that are somewhat separated in high dimensions to be far apart in 2-D is expensive. The t -distribution (with one degree of freedom, i.e. a Cauchy distribution) has much fatter tails, so the same KL gradient can accommodate larger spread without distorting the truly nearby points.

6.7.2 Examples

MNIST digits (6,000 points, 30-D PCA codes): t-SNE produces 10 well-separated clusters, one per digit class, with little overlap. Sammon's mapping, Isomap, and LLE all fail to achieve clean digit separation on the same data.

Olivetti faces: t-SNE groups images by identity into tight clusters; the other three methods produce much more mixed embeddings.

COIL-20 objects: each physical object, photographed under rotation, forms a closed loop (ring) in the 2-D t-SNE map — the topology of the rotation manifold is preserved. The other methods do not recover these loops clearly.

Iris data: PCA partially overlaps versicolor and virginica; t-SNE achieves perfect three-cluster separation with no overlap.

Multiple Tissue: both perplexity settings (30 and 50) give poor tissue separation — same conclusion as LLE and Isomap. The gene-expression data do not lie on a 2-D manifold.

6.8 Self-Organizing Maps

6.8.1 The Method

The **Self-Organizing Map** (SOM, also Kohonen map) pursues two goals simultaneously: **clustering** and **dimensionality reduction (down-projection)**.

A SOM maintains a lattice of K *neurons*. Each neuron k has a fixed grid position $\mathbf{y}_k \in \mathbb{R}^l$ (equidistantly filling a hypercube) and a learnable weight vector $\mathbf{w}_k \in \mathbb{R}^m$ in the input space. The goal is to preserve neighbourhood relations: points that are close in input space should activate neighbouring neurons (**topologically ordered maps**, TOMs).

Online update rule.. For each incoming sample $\mathbf{x}$:

1. Find the *winner* neuron: $k^* = \arg \max_s \mathbf{x}^\top \mathbf{w}_s$.
2. Move k^* and its neighbours towards $\mathbf{x}$:

$$\mathbf{w}_k^{\text{new}} = \mathbf{w}_k + \eta \delta(\|\mathbf{y}_k - \mathbf{y}_{k^*}\|) (\mathbf{x} - \mathbf{w}_k),$$

where η is the learning rate and $\delta(\cdot)$ is the *neighbourhood (window) function*: it equals 1 for $\mathbf{y}_k = \mathbf{y}_{k^*}$ and decreases monotonically with lattice distance.

SOM has no objective function

The SOM update rule cannot be derived from a single scalar energy or loss function. Practical consequences:

- No theoretical basis for choosing η , neighbourhood width, or stopping criterion.

- No guarantee that topographic ordering is achieved.
- No convergence proof.
- Different runs cannot be compared objectively.
- Overfitting is not detectable.
- The result carries no probability density and cannot be handed to a downstream probabilistic model.

GTM (Section 6.6) was developed as a principled probabilistic replacement.

6.8.2 Examples

Example 1 — 1-D SOM in a 2-D circle. Initially all weights are collapsed to a single point. Over $\sim 150,000$ updates the 1-D chain of neurons unfolds to trace the boundary of the circle, filling the circular data distribution.

Example 2 — 2-D SOM on a uniform square. A 10×10 grid of neurons progressively fills the square data space; by 300,000 steps the grid covers the entire region uniformly. *With a different random initialisation* the grid develops kinks that do **not** disappear even after 400,000 steps — a local minimum with no way to detect or escape it.

Non-uniform sampling. When data density is higher at the centre than the border, the SOM allocates more neurons to the centre — the grid is compressed there and stretched at the boundary, reflecting the input distribution.

Clustering

7.1 Mixture Models

7.1.1 The Method

A **mixture model** defines a generative distribution over observed data $\mathbf{x}$ by introducing a discrete latent variable $j \in \{1, \dots, l\}$ that indexes the component:

$$p(\mathbf{x} \mid \theta) = \sum_{j=1}^l p(j) p(\mathbf{x} \mid j, \theta_j),$$

where $p(j) = w_j \geq 0$ with $\sum_j w_j = 1$ are the *mixture weights* and $p(\mathbf{x} \mid j, \theta_j)$ is the j -th component density. The model provides a flexible way to represent multi-modal or multi-cluster distributions as a convex combination of simpler component distributions.

Bayesian cluster assignment. Given an observation $\mathbf{x}$, the posterior probability of belonging to component j is obtained via Bayes' theorem:

$$p(j \mid \mathbf{x}, \theta) = \frac{w_j p(\mathbf{x} \mid j, \theta_j)}{\sum_{t=1}^l w_t p(\mathbf{x} \mid t, \theta_t)}.$$

This *soft* assignment (also called *responsibility*) is the key quantity in the E-step of EM.

Log-likelihood gradient. The gradient of the log-likelihood with respect to component parameters θ_j factorises cleanly:

$$\frac{\partial}{\partial \theta_j} \log p(\mathbf{x} \mid \theta) = p(j \mid \mathbf{x}, \theta) \frac{\partial}{\partial \theta_j} \log p(\mathbf{x} \mid j, \theta_j).$$

The overall gradient is therefore the posterior-weighted sum of component gradients — exactly the structure exploited by EM.

7.1.2 Mixture of Gaussians

In a **mixture of Gaussians** (MoG) each component is a multivariate Gaussian $\mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$. EM iterates between computing soft assignments (E-step) and maximising the expected complete-data log-likelihood (M-step).

E-step. Compute the responsibility $r_{ij} = p(j \mid \mathbf{x}_i, \theta)$ for every sample–component pair:

$$r_{ij} = \frac{w_j \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}{\sum_{t=1}^l w_t \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_t, \boldsymbol{\Sigma}_t)}.$$

M-step. Update the parameters using the responsibilities:

$$w_j^{\text{new}} = \frac{1}{n} \sum_{i=1}^n r_{ij},$$

$$\boldsymbol{\mu}_j^{\text{new}} = \frac{\sum_i r_{ij} \mathbf{x}_i}{\sum_i r_{ij}},$$

$$\boldsymbol{\Sigma}_j^{\text{new}} = \frac{\sum_i r_{ij} (\mathbf{x}_i - \boldsymbol{\mu}_j^{\text{new}})(\mathbf{x}_i - \boldsymbol{\mu}_j^{\text{new}})^\top}{\sum_i r_{ij}}.$$

MAP-MoG. The ML solution can degenerate (a single-sample component with $\boldsymbol{\Sigma}_j \rightarrow \mathbf{0}$). Adding conjugate priors regularises the solution:

- **Wishart** prior on the precision: $\mathcal{W}(\boldsymbol{\Sigma}_j^{-1} \mid \alpha, \boldsymbol{\Psi})$.
- **Dirichlet** prior on the weights: $\mathcal{D}(\mathbf{w} \mid \gamma)$.
- **Gaussian** prior on the means: $\mathcal{N}(\boldsymbol{\mu}_j \mid \boldsymbol{\nu}_j, \eta^{-1} \boldsymbol{\Sigma}_j)$.

The MAP-EM updates replace ML formulas with:

$$w_j^{\text{new}} = \frac{\sum_i r_{ij} + \gamma - 1}{n + l(\gamma - 1)},$$

$$\boldsymbol{\mu}_j^{\text{new}} = \frac{\sum_i r_{ij} \mathbf{x}_i + \eta \boldsymbol{\nu}_j}{\sum_i r_{ij} + \eta},$$

$$\boldsymbol{\Sigma}_j^{\text{new}} = \frac{\sum_i r_{ij} (\mathbf{x}_i - \boldsymbol{\mu}_j^{\text{new}})(\mathbf{x}_i - \boldsymbol{\mu}_j^{\text{new}})^\top + 2\boldsymbol{\Psi} + \eta(\boldsymbol{\mu}_j^{\text{new}} - \boldsymbol{\nu}_j)(\boldsymbol{\mu}_j^{\text{new}} - \boldsymbol{\nu}_j)^\top}{\sum_i r_{ij} + 2\alpha + m + 2}.$$

Sensible defaults: $\alpha = m/2$, $\boldsymbol{\Psi} = \frac{1}{2}\mathbf{I}$ (or $\frac{1}{2}\text{cov}(\mathbf{x})$), $\gamma = 1$ (uniform Dirichlet, no prior on weights), $\eta = 0$ (no mean prior), $\boldsymbol{\nu}_j = \text{mean}(\mathbf{x})$.

Deriving the MAP weight update

Maximise $\log p(\mathbf{x} \mid \boldsymbol{\theta}) + \log p(\mathbf{w})$ subject to $\sum_j w_j = 1$ using a Lagrange multiplier λ :

$$\frac{\partial}{\partial w_j} \left[\sum_j \sum_i r_{ij} \log w_j + (\gamma - 1) \log w_j + \lambda(1 - \sum_j w_j) \right] = 0 \Rightarrow w_j = \frac{\sum_i r_{ij} + \gamma - 1}{\lambda}.$$

Summing over j and using $\sum_j w_j = 1$ gives $\lambda = n + l(\gamma - 1)$.

7.1.3 Mixture of Poissons

When data are **non-negative integer counts** the Gaussian assumption is inappropriate. A **mixture of Poissons** uses Poisson component distributions:

$$p(x) = \sum_{j=1}^l w_j \mathcal{P}(x; \lambda_j), \quad \mathcal{P}(x; \lambda_j) = \frac{\lambda_j^x e^{-\lambda_j}}{x!}.$$

The mean and variance of component j are both λ_j , so the model can capture over-dispersed count distributions via a heterogeneous mixture.

EM via Jensen’s inequality. The log-likelihood $\log \sum_j w_j \mathcal{P}(x_i; \lambda_j)$ contains a log of a sum, which is intractable to optimise directly. Introduce auxiliary variables $w_{ji} \geq 0$ with $\sum_j w_{ji} = 1$ and apply Jensen’s inequality to the concave log:

$$\log \sum_j w_j \mathcal{P}(x_i; \lambda_j) = \log \sum_j w_{ji} \cdot \frac{w_j \mathcal{P}(x_i; \lambda_j)}{w_{ji}} \geq \sum_j w_{ji} \log \frac{w_j \mathcal{P}(x_i; \lambda_j)}{w_{ji}}.$$

The bound is tight when $w_{ji} \propto w_j \mathcal{P}(x_i; \lambda_j)$, which defines the E-step.

E-step. Set

$$w_{ji}^{\text{new}} = \frac{w_j^{\text{old}} \mathcal{P}(x_i; \lambda_j^{\text{old}})}{p(x_i; \mathbf{w}^{\text{old}}, \boldsymbol{\lambda}^{\text{old}})}.$$

M-step. Maximise the lower bound over $\mathbf{w}$ and $\boldsymbol{\lambda}$:

$$w_j^{\text{new}} = \frac{1}{n} \sum_i w_{ji}, \quad \lambda_j^{\text{new}} = \frac{\sum_i w_{ji} x_i}{\sum_i w_{ji}}.$$

The λ_j update is the responsibility-weighted sample mean — exactly as expected for the sufficient statistic of the Poisson distribution.

Regularisation: a Dirichlet prior $\mathcal{D}(\mathbf{w} \mid \gamma)$ on the weights and a (non-informative) uniform prior on each λ_j .

7.1.4 Examples

Initialization strategies for MoG.. EM is sensitive to initialisation. Three strategies are common:

- **“em”**: run several short, low-tolerance EM passes from different random starts, then continue with a full precise run from the best solution found.
- **“rnd”**: purely random initialisations; pick the run with the highest final log-likelihood.
- **“svd”**: use the singular-value decomposition of the data matrix to derive a structured, data-adaptive starting point.

On a 2D synthetic dataset with 8 elongated clusters, the “em” and “svd” strategies reliably find the correct 8-cluster partition; “rnd” is more prone to poor local optima.

Covariance constraints.. The shape, volume, and orientation of each Gaussian component can be constrained to regularise the model or to reflect prior knowledge. Using Mclust notation, a three-letter code describes the multivariate constraint:

- **EII**: spherical, equal volume across components.
- **VII**: spherical, unequal volume.
- **EEI** / **VEI** / **EVI** / **VVI**: diagonal covariance (axis-aligned clusters), with equal/varying combinations of volume and shape.
- **EEE** / **EEV** / **VEV** / **VVV**: full ellipsoidal covariance with equal/varying volume, shape, and orientation.

Applied to a synthetic 2D dataset with 3 true clusters (banana-shaped, curved), a fully unconstrained MoG (VVV, 3 components) correctly recovers the three cluster shapes, while spherical or diagonal constraints produce boundary errors. Increasing to 6 components with tight covariance constraints causes over-splitting: each true cluster is divided into 2 spurious sub-clusters.

Multiple Tissue dataset.. MoG is applied to the Multiple Tissue dataset using the 76 genes with largest variance, projected onto the first two principal components. With 3 components (spherical, unequal volume) the three tissue types (liver, brain, adipose) are partially separated but brain and adipose overlap. With 4 components (spherical or diagonal) the partition improves; diagonal covariance with varying volume and shape gives a qualitatively better result, tracking the elongated liver cluster more accurately. Increasing to 5, 6, or 7 spherical components over-splits the clusters, showing that model selection (choosing l) is crucial.

7.2 k -Means Clustering

7.2.1 The Method

k -means is obtained from a mixture of Gaussians by imposing three simplifying assumptions:

1. equal weights $w_j = 1/l$;
2. spherical and equal covariance $\Sigma_j^{-1} = \mathbf{I}$;
3. *hard* (discrete) cluster membership — a sample belongs to exactly one cluster.

Under these assumptions the soft responsibility r_{ij} collapses to a binary indicator: $p(j \mid \mathbf{x}_i, \boldsymbol{\mu}_j) = 1$ if $j = \arg \min_k \|\mathbf{x}_i - \boldsymbol{\mu}_k\|$, zero otherwise. The only free parameters are the l cluster centres $\boldsymbol{\mu}_j$.

k -means algorithm

Initialise: choose l cluster centres $\boldsymbol{\mu}_j$.

Iterate until convergence:

1. *Assignment step*: assign each $\mathbf{x}_i$ to the nearest centre, $c_{\mathbf{x}_i} = \arg \min_k \|\mathbf{x}_i - \boldsymbol{\mu}_k\|$.
2. *Update step*: recompute each centre as the mean of its assigned points, $\boldsymbol{\mu}_j^{\text{new}} = \frac{1}{n_j} \sum_{i: c_{\mathbf{x}_i}=j} \mathbf{x}_i$, where $n_j = \#\{i : c_{\mathbf{x}_i} = j\}$.

Properties: k -means is fast and robust to outliers (hard assignment prevents a single outlier from shifting all means), but it is simple (only spherical, equal-size clusters) and sensitive to initialisation. A centre placed near outliers can remain stuck there, failing to model the true cluster distribution.

Fuzzy k -means replaces the hard membership with a soft (“softmax”) membership controlled

by a parameter $b > 1$:

$$p^b(j \mid \mathbf{x}_i, \boldsymbol{\mu}_j) = \frac{\|\mathbf{x}_i - \boldsymbol{\mu}_j\|^{-2/(b-1)}}{\sum_{k=1}^l \|\mathbf{x}_i - \boldsymbol{\mu}_k\|^{-2/(b-1)}}.$$

The centre update becomes the responsibility-weighted mean:

$$\boldsymbol{\mu}_j^{\text{new}} = \frac{\sum_i p^b(j \mid \mathbf{x}_i, \boldsymbol{\mu}_j) \mathbf{x}_i}{\sum_i p^b(j \mid \mathbf{x}_i, \boldsymbol{\mu}_j)},$$

and the minimised objective is $\sum_j \sum_i p^b(j \mid \mathbf{x}_i, \boldsymbol{\mu}_j) \|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2$. As $b \rightarrow 1$ this recovers hard k -means; as $b \rightarrow \infty$ all memberships converge to $1/l$ (fully fuzzy).

7.2.2 Examples

Synthetic 2D data (5 clusters). On a synthetic dataset with 5 Gaussian clusters, k -means with $l = 5$ recovers the true partition (cluster centres at the true means) when initialised well. Local minima arise frequently: one failure mode is that one model centre explains two true clusters while a second true cluster is divided between two model centres; another is that three model centres all compete for the same true cluster while one true cluster is unmodelled. These failures highlight why multiple random restarts are necessary in practice.

Synthetic 2D data ($k = 8$). With $l = 8$ and 5 true clusters the model must split some clusters. Different initialisations produce markedly different partitions, confirming the strong dependence on initialisation.

Iris data. k -means with $l = 3$ on the first two PCA components of the Iris dataset typically recovers two of the three species exactly, with errors only at the border between the two overlapping classes (*Iris versicolor*/*Iris virginica*). Occasionally a suboptimal initialisation merges the two overlapping clusters and splits the well-separated one — these solutions have higher objective values but cannot always be distinguished from the correct solution by the k -means objective alone.

Multiple Tissue data. k -means with $l = 4$ on the Multiple Tissue dataset correctly identifies the four tissue types in the most typical run. Suboptimal initialisations place a centre in the empty region between the adipose and brain clusters, producing a qualitatively different partition.

7.3 Hierarchical Clustering

7.3.1 The Method

Hierarchical clustering represents the cluster structure as a *dendrogram* — a tree whose leaves are individual samples and whose internal nodes encode the distance at which two sub-trees were merged (or split). Cutting the dendrogram at a chosen height yields a flat partition; varying the cut gives a nested family of solutions at every granularity.

Two broad strategies exist:

- **Agglomerative (bottom-up):** start with n singleton clusters and greedily merge the closest pair at each step.
- **Divisive (top-down):** start with one cluster containing all points. A common approach builds the *minimal spanning tree* of the pairwise distance graph and then removes the largest edge to split into two clusters, repeating recursively. Alternatively, remove the edge that is considerably larger than other edges within its cluster.

Linkage criteria define the distance between two clusters A and B :

$$\begin{aligned}
 d_{\min}(A, B) &= \min_{a \in A, b \in B} \|a - b\| && \text{(single linkage)} \\
 d_{\max}(A, B) &= \max_{a \in A, b \in B} \|a - b\| && \text{(complete linkage)} \\
 d_{\text{avg}}(A, B) &= \frac{1}{n_A n_B} \sum_{a \in A} \sum_{b \in B} \|a - b\| && \text{(average linkage / UPGMA)} \\
 d_{\text{mean}}(A, B) &= \|\bar{a} - \bar{b}\| && \text{(centroid / mean linkage)}
 \end{aligned}$$

where n_A, n_B are the cluster sizes and $\bar{a}, \bar{b}$ are the cluster means. The element-level distance $\|\cdot\|$ can be Euclidean, Manhattan, or Mahalanobis.

Linkage behaviour

Complete linkage avoids elongated clusters — it guarantees that the diameter (largest intra-cluster distance) stays small — but clusters may not be well-separated. **Single linkage** ensures any two points from different clusters are far apart; it is used for *leave-one-cluster-out* cross-validation where an entire new group is withheld. **Average linkage (UPGMA)** is a commonly used compromise.

Ward's method merges the pair of clusters whose union minimises the total within-cluster sum of squares — a variance-based criterion that tends to produce compact, spherical clusters of similar size.

7.3.2 Examples

US Arrest data.. The 50 US states are clustered on four crime-rate variables. Ward's dendrogram produces two clean super-clusters (high-crime vs. low-crime states) with a clear gap, giving an unambiguous 2-cluster cut. Single linkage exhibits *chaining*: states are added one at a time to a growing chain, producing an unbalanced dendrogram with little internal structure. Complete and average linkage produce intermediate results, broadly consistent with Ward's partition. McQuitty, median, and centroid methods give similar dendrograms to average linkage on this dataset.

Synthetic 5-cluster data.. Ward's distance produces a perfect recovery of all 5 Gaussian clusters when the dendrogram is cut at the right height. Complete and average linkage also recover the correct partition. Single linkage again suffers from chaining and produces a poor cut.

Iris data.. With 3 clusters, Ward's and average linkage both achieve good (but not perfect) separations — a few *Iris versicolor*/*Iris virginica* points are misassigned at the class boundary. Single linkage chains through the overlapping region and fails to isolate the third species cleanly. Increasing to 5 clusters over-splits the data; the extra clusters appear at boundary points marked by misassignment arrows.

Multiple Tissue data.. With 4 clusters, Ward's linkage correctly groups the four tissue types (liver, brain, adipose, colon/kidney mix) with only a few misassigned points. Complete and average linkage with 4 clusters produce similar quality. Single linkage with 4 clusters chains the liver samples into one huge cluster and isolates small outliers instead of identifying the biological tissue groups. Increasing to 6 clusters with Ward's or average linkage over-splits the data, creating biologically meaningless sub-clusters.

7.4 Similarity-Based Clustering

Similarity-based clustering groups objects using only pairwise similarities, without requiring an explicit feature-vector representation. Similarities can come from web links, social interactions, co-occurrence counts (co-expressed genes, co-cited papers), spatial distances, sequence alignments, or structural comparisons.

7.4.1 Aspect Model

The **aspect model** applies to discrete data of co-occurring pairs (x, y) — for example, person x buys product y , or word x appears in document y . The data are counts $n(x, y)$.

Aspect model

Introduce a hidden class variable $z \in \{z_1, \dots, z_l\}$. Conditional on z , x and y are independent:

$$p(x, y) = \sum_z p(z) p(x | z) p(y | z).$$

The hidden factor z acts as the common cause of both x and y , explaining their correlation without direct dependence. *Example:* hot-chocolate consumption and crime rate are both driven by cold weather (z), so $x \perp y | z$.

EM for the aspect model.

- **E-step:** compute the posterior $p(z | x, y) = \frac{p(z) p(x | z) p(y | z)}{\sum_{z'} p(z') p(x | z') p(y | z')}$.
- **M-step:** update the parameters using the weighted counts $n(x, y)$:

$$\begin{aligned} p(z) &= \sum_{x,y} n(x, y) p(z | x, y), \\ p(y | z) &= \frac{\sum_x n(x, y) p(z | x, y)}{p(z)}, \\ p(x | z) &= \frac{\sum_y n(x, y) p(z | x, y)}{p(z)}. \end{aligned}$$

Clustering. Objects x are clustered via the marginal posterior $p(z | x) \propto p(x | z) p(z)$, where $p(x) = \sum_{z,y} p(z) p(x | z) p(y | z)$. Each value of z represents one cluster; analogous formulas cluster y or pairs (x, y) .

7.4.2 Affinity Propagation

The Method

Affinity propagation (AP) is a similarity-based, *exemplar-based* clustering method: cluster centres must be actual data points (*exemplars* or *prototypes*). Unlike k -centres clustering — which also enforces this constraint but requires a good initialisation and works only for a small number of clusters — AP determines both the exemplars and the cluster memberships simultaneously via message passing, without needing the number of clusters specified in advance.

Four quantities drive the algorithm:

- **Similarities** $s(i, k)$: how well object k can serve as an exemplar for object i (e.g. negative squared distance, $s(i, k) = -\|\mathbf{x}_i - \mathbf{x}_k\|^2$).

- **Preferences** $s(k, k)$: the prior probability that object k becomes an exemplar; a uniform value controls the expected number of clusters. *Higher preference $\rightarrow$ more cluster centres; lower preference $\rightarrow$ fewer cluster centres.*
- **Responsibilities** $r(i, k)$: message from i to candidate exemplar k , reflecting how strongly k should be the exemplar for i relative to competing candidates.
- **Availabilities** $a(i, k)$: message from candidate exemplar k to object i , reflecting how much evidence k collects from other objects that it should be an exemplar.

Affinity Propagation update rules

Initialise $a(i, k) = 0$. Iterate, applying damping with factor $\lambda \in [0, 1)$ to improve numerical stability (e.g. $\lambda = 0.5$ by default):

$$r(i, k) = s(i, k) - \max_{k', k' \neq k} \{a(i, k') + s(i, k')\},$$

$$a(i, k) = \min\left\{0, r(k, k) + \sum_{i', i' \notin \{i, k\}} \max\{0, r(i', k)\}\right\} \quad (i \neq k),$$

$$a(k, k) = \sum_{i', i' \neq k} \max\{0, r(i', k)\}.$$

Object i is assigned to the exemplar $k^* = \arg \max_k \{a(i, k) + r(i, k)\}$.

Each message is updated as a weighted average of the new value and the previous value: $r_{\text{new}} \leftarrow (1 - \lambda) r_{\text{computed}} + \lambda r_{\text{old}}$ (same for a). A larger λ makes convergence slower but more stable; a smaller λ converges faster but may oscillate.

Intuition behind the messages

The responsibility $r(i, k)$ drops when a competing candidate k' has high similarity and availability to i — it measures how uniquely suited k is to represent i . The self-availability $a(k, k)$ grows when many other objects send positive responsibilities to k , confirming k as a strong exemplar. Together the two messages implement a competition: candidates that gather broad support become exemplars; others do not.

Limitation. AP implicitly assumes spherical clusters of equal size because similarities are typically Euclidean distances. When cluster sizes differ, AP cannot adjust the effective bandwidth per cluster, and elements of a large diffuse cluster may be assigned to a nearby compact cluster instead.

Examples

Toy graph (convergence illustration).. Starting from a fully connected graph of ~ 25 nodes, the messages evolve over 6 iterations: edges redistribute, 3–4 exemplars emerge (highlighted in red), and by convergence each point points exclusively to its exemplar.

Face images.. On a face database AP selects a small number of representative faces (exemplars, highlighted in coloured boxes); all other faces point to their nearest exemplar. Compared with k -centres clustering run on the same data, AP finds higher-quality exemplars for the 15 hardest-to-represent images.

Airport clustering.. AP applied to air-travel efficiency between 456 North American airports identifies 7 hub exemplars. The assignment reflects domestic vs. international flight patterns — e.g.

the Canada–USA border separates Toronto and Philadelphia hubs. The dense Seattle–Vancouver corridor causes Pacific northwest cities to be assigned to the Seattle hub.

Synthetic 2D data: balanced clusters (x7).. On a dataset with 6 clusters (3 large / 3 small by variance), AP correctly identifies 6 exemplars and recovers the clean partition when all clusters are roughly spherical and well-separated.

Synthetic 2D data: mixed-size clusters (x7A, x9, d6).. When one dense small cluster sits inside or near a large diffuse cluster, AP merges them — it cannot detect the small cluster inside the large one because it assumes equal cluster variances. On the x9 dataset (6 clusters of 2 small, 2 medium, 2 large) AP conflates one pair, producing 5 clusters instead of 6. On the 6-D d6/d6A datasets (one large + three small clusters, 100 samples, PCA visualisation), AP again assigns elements of the large cluster to smaller neighbours — the large cluster has no single exemplar that would win the competition.

Iris data.. AP with default preference ($p = 0$) produces a good 3-cluster partition with a few misassignments at the *Iris versicolor*/*Iris virginica* boundary, similar to k -means and MoG.

Multiple Tissue data.. With preference $p = 0$ AP finds 4 clusters that largely match the tissue types, with a few misassigned samples. Increasing the preference to $p = 1$ over-clusters into many singleton exemplars (every point becomes its own cluster), showing sensitivity to the preference hyperparameter.

7.4.3 Similarity-based Mixture Models

Similarity-based Mixture of Gaussians

Similarity-based mixture models are mixture models that operate solely on pairwise similarities $k(\mathbf{x}_j, \mathbf{x}_i)$ between the n training points — no explicit feature vectors are used. This is closely related to kernel methods: the input to the algorithm is the $n \times n$ kernel matrix $K_{ij} = k(\mathbf{x}_j, \mathbf{x}_i)$.

Starting from a Gaussian mixture $p(\mathbf{x}) = \sum_{j=1}^K p(j) k(\mathbf{x}; \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$, introduce a scalar variance σ_i^2 per component and define

$$k(\mathbf{x}; \boldsymbol{\mu}_i, \sigma_i^2) = \frac{1}{\sigma_i^m} k(\mathbf{x}, \boldsymbol{\mu}_i)^{1/\sigma_i^2},$$

so the model becomes

$$p(\mathbf{x}) = \sum_{j=1}^K p(j) \frac{1}{\sigma_j^m} k(\mathbf{x}, \boldsymbol{\mu}_j)^{1/\sigma_j^2}.$$

The component variance σ_j^2 controls the width of cluster j , allowing the model to handle clusters of different sizes — a key advantage over affinity propagation.

Connection to standard MoG

If $k(\mathbf{x}, \boldsymbol{\mu}_i) = \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_i, \mathbf{I})$ and $\sigma_i^2 = 1$ for all i , this reduces to a standard spherical MoG. The similarity-based formulation generalises this to arbitrary positive semi-definite kernels, and expresses all update rules solely in terms of the given similarities $k(\mathbf{x}_k, \mathbf{x}_j)$.

Representing the Centers

Setting $\alpha_i = p(i)$ and $\hat{\alpha}_{ik} = p(i \mid \mathbf{x}_k, \boldsymbol{\alpha})$, the weights satisfy

$$\alpha_i = \frac{1}{n} \sum_{k=1}^n \hat{\alpha}_{ik}.$$

The EM centre update requires evaluating $\|\mathbf{x}_l - \boldsymbol{\mu}_i\|^2$. Assuming the Gaussian kernel $k(\mathbf{x}_k, \mathbf{x}_j) = \exp(-\frac{1}{2}\|\mathbf{x}_k - \mathbf{x}_j\|^2)$, so that $\|\mathbf{x}_k - \mathbf{x}_j\|^2 = -2 \log k(\mathbf{x}_k, \mathbf{x}_j)$, one can show:

$$\|\mathbf{x}_l - \boldsymbol{\mu}_i\|^2 = \frac{-2 \sum_k \hat{\alpha}_{ik} \log k(\mathbf{x}_k, \mathbf{x}_l)}{\sum_k \hat{\alpha}_{ik}} + \frac{\sum_{k,j} \hat{\alpha}_{ik} \hat{\alpha}_{ij} \log k(\mathbf{x}_k, \mathbf{x}_j)}{\left(\sum_k \hat{\alpha}_{ik}\right)^2}.$$

This expresses the squared distance to the (virtual) centre entirely through pairwise similarities — no explicit $\boldsymbol{\mu}_i$ needs to be stored or updated. The updated kernel value is $k(\mathbf{x}_i, \boldsymbol{\mu}_i)^{\text{new}} = \exp(-\frac{1}{2}\|\mathbf{x}_i - \boldsymbol{\mu}_i\|^2)$.

The full covariance $\boldsymbol{\Sigma}_i$ could also be expressed via similarities, but this requires inverting an $n \times n$ matrix, which is computationally expensive and numerically unstable; only the scalar variance σ_i^2 is therefore estimated.

Update Rules

Define $\hat{\alpha}_i = \frac{1}{n} \sum_k \hat{\alpha}_{ik}$ and $\gamma_s = \sum_i \gamma_i$. The complete EM update cycle is:

1. **E-step:** $\hat{\alpha}_{ik} = \frac{\alpha_i^{\text{old}} k(\mathbf{x}_k; \boldsymbol{\mu}_i, \sigma_i^2)}{p(\mathbf{x}_k)}$, $p(\mathbf{x}_k) = \sum_i \alpha_i^{\text{old}} k(\mathbf{x}_k; \boldsymbol{\mu}_i, \sigma_i^2)$.
2. **Weight update:** $\alpha_i^{\text{new}} = \frac{n\hat{\alpha}_i + \gamma_i - 1}{\gamma_s}$.
3. **Variance update:** $(\sigma_i^2)^{\text{new}} = \frac{\sum_k \hat{\alpha}_{ik} (-2 \log k(\mathbf{x}_k, \boldsymbol{\mu}_i)) + wS}{n\hat{\alpha}_i + w}$.
4. **Kernel update:** compute $-2 \log k(\mathbf{x}_l, \boldsymbol{\mu}_i) = \|\mathbf{x}_l - \boldsymbol{\mu}_i\|^2$ from the similarity formula above, then set $k(\mathbf{x}_l, \boldsymbol{\mu}_i)^{\text{new}} = \exp(-\frac{1}{2}\|\mathbf{x}_l - \boldsymbol{\mu}_i\|^2)$.

Hyperparameters: γ_i from a Dirichlet prior on $\boldsymbol{\alpha}$; w and S from a Wishart prior on the precisions $1/\sigma_i^2$.

Examples

x7 data (6 balanced clusters).. Both affinity propagation and EMcluster (similarity-based MoG EM) correctly recover all 6 clusters. The α_i and σ_i^2 traces over 200 EM iterations show rapid convergence: all α_i converge to $\approx 1/6$ and σ_i^2 differentiates the two cluster sizes (small clusters settle to lower variance).

x7A data (overlapping/mixed-size clusters).. AP merges one large and one small cluster; EMcluster correctly separates all 6.

x9 data (6 clusters: 2 small / 2 medium / 2 large).. AP merges the two closest clusters; EMcluster identifies all 6 correctly.

d6/d6A data (1 large + 3 small, 6-D).. AP cannot detect the large diffuse cluster and assigns its elements to smaller neighbours. EMcluster correctly identifies the large cluster by fitting a large σ^2 to it, while the small clusters get small σ^2 values.

The key advantage of the similarity-based MoG over AP is the per-component variance parameter σ_i^2 : because AP uses raw similarities without adapting cluster width, it implicitly assumes equal-size spherical clusters and fails when cluster sizes differ substantially.

Biclustering

8.1 Types of Biclusters

Biclustering simultaneously clusters both the rows and the columns of a data matrix. A *bicluster* is a pair (I, J) where $I \subseteq \{1, \dots, n\}$ is a subset of rows (e.g. genes) and $J \subseteq \{1, \dots, m\}$ is a subset of columns (e.g. samples), such that the entries X_{ij} for $i \in I, j \in J$ exhibit a coherent pattern. Unlike standard clustering, rows and columns can belong to multiple biclusters simultaneously, or to none at all.

The literature distinguishes six types of biclusters ordered by increasing structural complexity:

Types of Biclusters

Let a_{ij} denote the entry in row i , column j of the bicluster sub-matrix.

- **Constant values:** $a_{ij} = c$ for all i, j — all entries identical.
- **Constant rows:** $a_{ij} = r_i$ — each row has a fixed value across all bicluster columns.
- **Constant columns:** $a_{ij} = c_j$ — each column has a fixed value across all bicluster rows.
- **Additive coherent values:** $a_{ij} = \mu + r_i + c_j$ — row and column offsets are additive; rows are shifted versions of each other.
- **Multiplicative coherent values:** $a_{ij} = \mu \cdot r_i \cdot c_j$ — rows are scalar multiples of each other; columns are scalar multiples of each other.
- **General coherent values:** rows (or columns) are related by a more general function; includes the above as special cases.

Intuition & Mental Model

An additive coherent bicluster with a 5×5 sub-matrix can be visualised as: each entry equals a row offset r_i plus a column offset c_j plus a global mean μ . Subtracting row means from each row flattens all rows to the same profile, confirming the additive structure. Multiplicative coherence means all rows are proportional — dividing each row by its mean yields identical rows.

8.2 Overview of Biclustering Methods

Biclustering algorithms fall into four broad categories:

Variance minimization. These methods search for sub-matrices with low internal variance. *Cheng–Church δ -biclusters* define the mean squared residue of a bicluster as

$$H(I, J) = \frac{1}{|I||J|} \sum_{i \in I, j \in J} (a_{ij} - a_{iJ} - a_{IJ} + a_{IJ})^2,$$

and accept any bicluster with $H(I, J) \leq \delta$. Related methods include δ -cluster and δ -pClusters.

Two-way clustering. Iterative co-clustering procedures alternate between clustering rows and columns. Representative methods: CTWC (Coupled Two-Way Clustering), ITWC (Iterative Two-Way Clustering), and DCC (Double Conjugated Clustering).

Motif and pattern recognition. Search for biclusters defined by expression patterns or ordinal relationships. Methods include xMOTIF (expression motif), OPSM (order-preserving sub-matrices), SPEC (spectral biclustering), and ISA (Iterative Signature Algorithm).

Probabilistic and generative models. Fit a probabilistic model to the data and infer biclusters from the model parameters. Methods include SAMBA (probabilistic bipartite graph), PRMs, ProBic, cMonkey, Plaid model, BBC (Bayesian BiClustering), and **FABIA** (Factor Analysis for Bicluster Acquisition).

8.3 FABIA Biclustering

FABIA (Factor Analysis for Bicluster Acquisition) represents each bicluster as an *outer product* of two sparse vectors.

FABIA Bicluster Representation

A bicluster is the rank-1 matrix $\mathbf{u}_j \mathbf{y}_j^T$ where

- $\mathbf{u}_j \in \mathbb{R}^m$ is a sparse *prototype* (loading) vector — its non-zero entries indicate which features (genes) belong to bicluster j .
- $\mathbf{y}_j \in \mathbb{R}^n$ is a sparse *factor* vector — its non-zero entries indicate which samples belong to bicluster j .

The full data matrix is modelled as

$$X = \sum_{j=1}^l \mathbf{u}_j \mathbf{y}_j^T + \Gamma = YU^T + \Gamma,$$

where $X, \Gamma \in \mathbb{R}^{n \times m}$, $U \in \mathbb{R}^{m \times l}$, $Y \in \mathbb{R}^{n \times l}$, and Γ is noise. Per sample: $\mathbf{x}_i = U \tilde{\mathbf{y}}_i + \boldsymbol{\varepsilon}_i$.

FABIA as Factor Analysis

The per-sample model $\mathbf{x} = U \tilde{\mathbf{y}} + \boldsymbol{\varepsilon}$ is a standard factor analysis (FA) model:

- $U \in \mathbb{R}^{m \times l}$ is the loading matrix (factor loadings).
- $\tilde{\mathbf{y}} \sim \mathcal{N}(\mathbf{0}, I)$ in standard FA.
- $\boldsymbol{\varepsilon} \sim \mathcal{N}(\mathbf{0}, \Psi)$ with Ψ diagonal (independent noise per feature).
- U explains the *common variance*; Ψ explains the *independent variance*.

Sparseness via Laplace Priors

Standard FA does not produce sparse solutions. FABIA enforces sparseness by replacing the Gaussian prior on factors with a **Laplace prior**:

FABIA Laplace Priors

$$p(\tilde{\mathbf{y}}) = \left(\frac{1}{\sqrt{2}}\right)^l \prod_{j=1}^l \exp(-\sqrt{2} |\tilde{y}_j|),$$

$$p(\mathbf{u}_j) = \left(\frac{1}{\sqrt{2}}\right)^m \prod_{k=1}^m \exp(-\sqrt{2} |u_{kj}|).$$

These are Laplace (double-exponential) distributions, which have a sharp peak at zero, promoting sparse solutions.

Because the posterior under Laplace priors is intractable, FABIA uses **variational EM**: a variational approximation to the posterior is maximised in the E-step, and the model parameters are updated in the M-step.

Intuition & Mental Model

The Laplace prior is a sparsity-inducing prior — analogous to the ℓ_1 penalty in LASSO. Factors and loadings are driven towards zero unless there is strong evidence in the data, so only the samples and features that genuinely participate in a bicluster receive large values. The resulting sparse outer products $\mathbf{u}_j \mathbf{y}_j^T$ cleanly delimit rectangular bicluster regions in the data matrix.

8.4 Examples

Synthetic data. On a 1000×100 dataset with 13 planted biclusters, FABIA recovers the correct bicluster structure: the absolute factor matrix $|Y|$ (samples $\times$ biclusters) and the absolute loading matrix $|U|$ (features $\times$ biclusters) both show clear block structure aligned with the planted biclusters. The reconstructed data $\hat{X} = YU^T$ visually reproduces the block pattern in the original matrix.

Super-Gaussian 50-dimensional data. FABIA recovers latent structure in simulated super-Gaussian data with 50 features, demonstrating robustness beyond the Gaussian assumption of standard FA.

Iris dataset. FABIA identifies two informative biclusters:

- *Bicluster 2* — petal features (petal length, petal width); separates *Iris setosa* from the other two species.
- *Bicluster 3* — sepal features (sepal length, sepal width); provides further separation among *Iris versicolor* and *Iris virginica*.

A biplot visualises both features and samples in the same 2D space defined by the two biclusters.

Multiple Tissue dataset. FABIA discovers tissue-specific expression signatures, with quality improving with more iterations:

- *500 iterations* — bicluster 2 separates prostate (green); bicluster 3 separates colon (orange); a fourth bicluster separates lung (blue).

- *2000 iterations* — bicluster 4 clearly separates prostate; bicluster 1 separates colon; biclusters 3 and 2 jointly help separate lung (blue) and breast (red). More iterations produce sparser and more clearly separated solutions.

Because the FABIA solution is sparse, the top genes of a bicluster are directly interpretable. For the prostate bicluster the most relevant genes are KLK3, ACPP, KLK2, CUTL2, and RNF41 — all strongly associated with prostate tissue: KLK3 is the *prostate-specific antigen* (PSA), ACPP is secreted by the epithelial cells of the prostate gland, and KLK2 cleaves pro-PSA into its active form. The biplot (biplots of bicluster 1 vs. 2 and of bicluster 1 vs. 3) confirms this: large circles mark the features (genes) that drive the bicluster; sparseness pushes most other features to zero, collapsing them onto the axes, so the signal is clean.

Breast cancer dataset. Both PCA and FABIA biclustering separate the blue subclass clearly, but both have difficulty separating the red subclass. The biplot of biclusters 2 and 3 shows the three subclasses (cyan, green, blue) well separated with distinctive marker genes (PARBL, COL14A1, CRYAB, MIA visible in the biplot).

DLBCL dataset. PCA separates the blue subclass; FABIA separates subclasses better because the red cluster is more clearly isolated in bicluster space than in principal component space. The biplot of biclusters 1 and 2 shows good subclass separation with NME1 as a prominent marker gene.

Hidden Markov Models

9.1 Hidden Markov Models in Bioinformatics

Hidden Markov models are the most prominent *generative model* in bioinformatics. The generative perspective means that the model defines a probability distribution over output sequences; new sequences can be scored by their likelihood and sequences can be generated that are statistically similar to the training sequences. HMMs are well suited to protein and DNA sequences, which are discrete.

Gene prediction. HMMs identify exons and introns along a genomic sequence. The tool GENSCAN achieves a base-pair specificity of 50%–80%.

Profile HMMs. A profile HMM stores a *multiple sequence alignment* in a model; new sequences are scored by their likelihood under the profile. When building a profile HMM from unaligned sequences, training can get stuck in local maxima; *deterministic annealing* (used in HMMER) avoids this. Standard software packages: HMMER and SAM. The PFAM protein family database is built entirely on profile HMMs.

Profile HMMs have inherent limitations:

- Cannot discover long-range dependencies.
- Cannot handle real-valued inputs (discrete alphabet only).
- Cannot detect higher-order correlations between positions.
- Cannot use a negative training set (only model what sequences look like, not what they are not).

Other applications. Remote homology detection, sequence scoring, HMMs combined with phylogenetic trees, and family databases (PFAM).

9.2 Hidden Markov Model Basics

An HMM is a graph of connected hidden states. The hidden state at time t (sequence position t) is $u_t \in \{1, \dots, S\}$. Each state produces an observation according to a probabilistic output distribution. The model evolves by jumping to another state or staying in the current state at each step. The set of all possible state sequences can be visualised as a trellis: S rows (states) and T columns (time steps), with each left-to-right path corresponding to one possible hidden sequence.

HMM Parameters

- **Start probability:** $p_S(u_1)$ — probability of being in state u_1 at position 1.
- **Transition probability:** $p(u_t | u_{t-1})$ — probability of transitioning from state u_{t-1} to state u_t .
- **Emission probability:** $p(x_t | u_t)$ — probability of emitting observation x_t from state u_t .

Markov assumption: the next state depends only on the current state (first-order). Higher-order HMMs condition on more history, e.g. second-order: $p(u_t | u_{t-1}, u_{t-2})$.

The probability of a hidden state sequence $u^T = (u_1, u_2, \dots, u_T)$ of length T is:

$$p(u^T) = p_S(u_1) \prod_{t=2}^T p(u_t | u_{t-1}). \quad (9.1)$$

The joint probability of a state sequence and an observation sequence $x^T = (x_1, \dots, x_T)$ is:

$$p(x^T, u^T) = p_S(u_1) p(x_1 | u_1) \prod_{t=2}^T p(u_t | u_{t-1}) p(x_t | u_t). \quad (9.2)$$

Intuition & Mental Model

A simple two-state HMM has states *on* ($u = 1$) and *off* ($u = 0$) with self-loops (staying in the current state) and cross-transitions. Each state emits a characteristic output distribution. The trellis picture makes clear that computing $p(x^T)$ by summing over all S^T paths is exponential — the forward algorithm provides an efficient dynamic-programming solution.

9.3 Expectation Maximization for HMM: Baum-Welch Algorithm**9.4 Viterbi Algorithm****9.5 Input Output Hidden Markov Models****9.6 Factorial Hidden Markov Models****9.7 Memory Input Output Factorial Hidden Markov Models****9.8 Tricks of the Trade****9.9 Profile Hidden Markov Models**

Boltzmann Machines

10.1 The Boltzmann Machine

10.2 Learning in the Boltzmann Machine

10.3 The Restricted Boltzmann Machine
